Name:S.Venkata Akshay
Reg no: 11259A399

A. Creating Arrays

Aim

To write a Python program using the NumPy library to create and display different types of arrays, such as 1D arrays, 2D arrays, zeros, ones, identity matrices, and range of numbers.

Algorithm

1.Start the program.

2.Import the NumPy library as np.

3.Create a 1D array using np.array().

4.Create a 2D array using np.array().

5.Create a zero matrix using np.zeros().

6.Create a matrix of ones using np.ones().

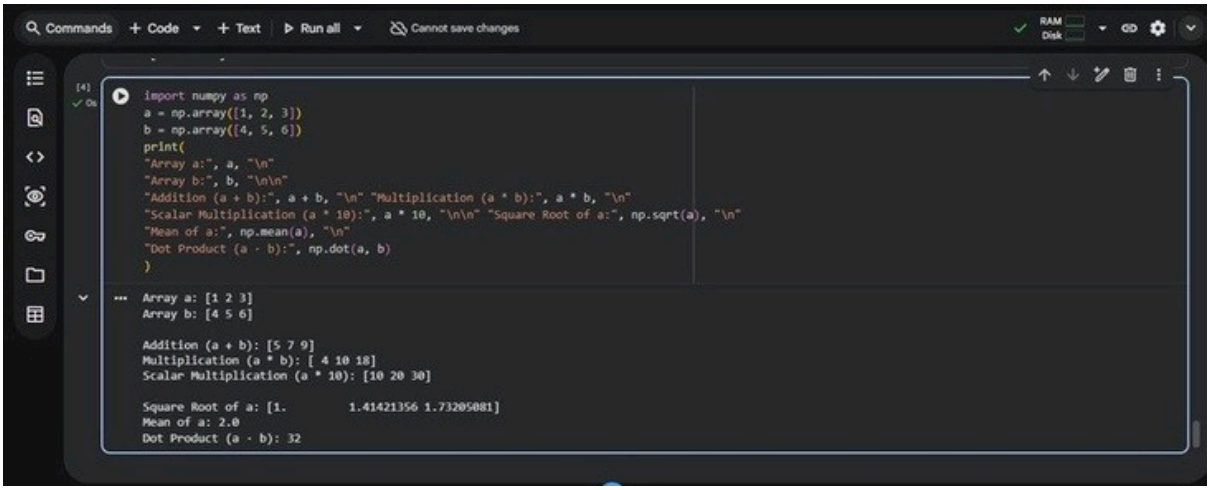
7.Create an identity matrix using np.eye().

8.Generate a sequence of numbers using np.arange().

9.Display all the arrays using the print() function.

10. Stop the program.

Your paragraph text



Result

Thus, the Python program to create and display different types of NumPy arrays was successfully executed.

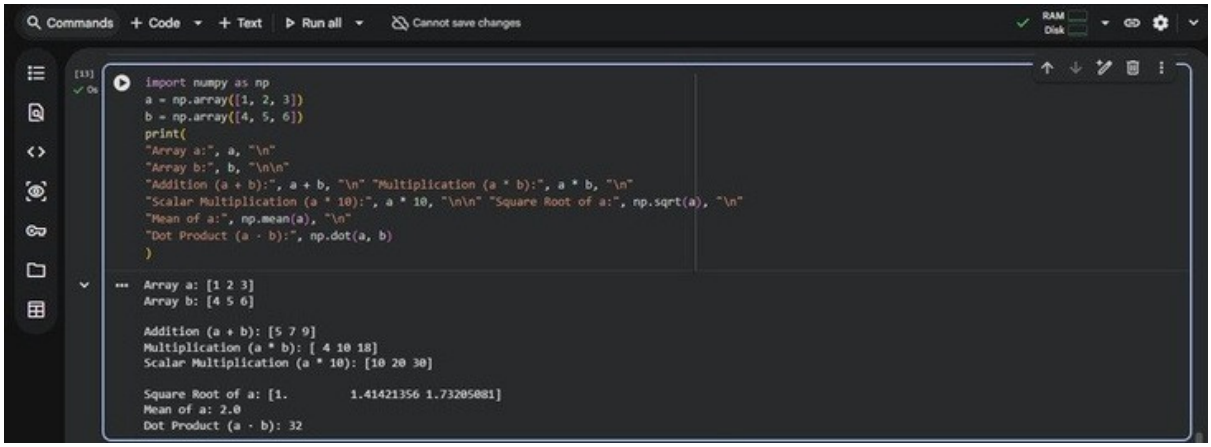
B. Array Operations

Aim

To write a Python program using NumPy to perform array operations such as addition, multiplication, scalar multiplication, square root, mean, and dot product.

Algorithm

1. Start the program.
2. Import the NumPy library as np.
3. Create two arrays, a = [1, 2, 3] and b = [4, 5, 6].
4. Display both arrays.
5. Perform element-wise addition using a + b.
6. Perform element-wise multiplication using a * b.
7. Perform scalar multiplication using a * 10.
8. Calculate the square root of array a using np.sqrt().
9. Calculate the mean of array a using np.mean().
10. Calculate the dot product of arrays a and b using np.dot().
11. Display all the results.
12. Stop the program.



Result

Thus, the Python program to perform various NumPy array operations, including addition, multiplication, scalar multiplication, square root, mean, and dot product, was successfully executed.

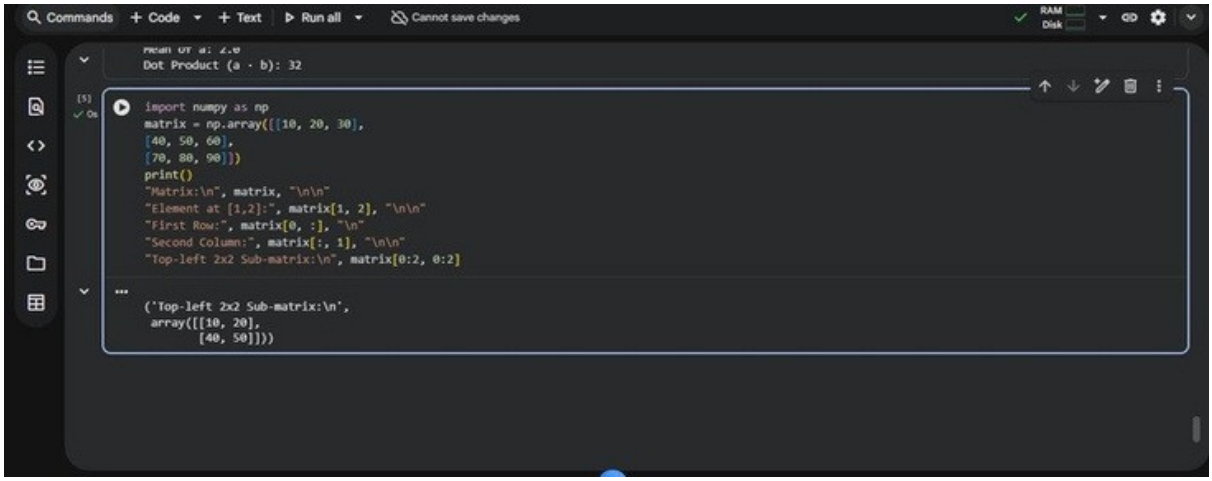
C. Indexing and Slicing

Aim

To write a Python program using NumPy to access individual elements, extract rows and columns, and obtain a sub-matrix from a two-dimensional array using indexing and slicing.

Algorithm

1. Start the program.
2. Import the NumPy library as np.
3. Create a 3×3 matrix containing the elements 10 to 90.
4. Display the matrix.
5. Access the element at index [1, 2] using matrix[1, 2].
6. Extract the first row using matrix[0, :].
7. Extract the second column using matrix[:, 1].
8. Extract the top-left 2×2 sub-matrix using matrix[0:2, 0:2].
9. Display all the extracted elements and sub-matrices.
10. Stop the program.



Result

Thus, the Python program to perform indexing and slicing operations on a NumPy matrix was successfully executed, and the required elements, row, column, and sub-matrix were obtained.

A. Creating DataFrames

Aim

To write a Python program using the Pandas library to create a DataFrame from a dictionary and display the first few rows, data types, non-null values, and statistical summary.

Algorithm

1.Start the program.

2.Import the Pandas library as pd.

3.Create a dictionary containing Name, Age, and City.

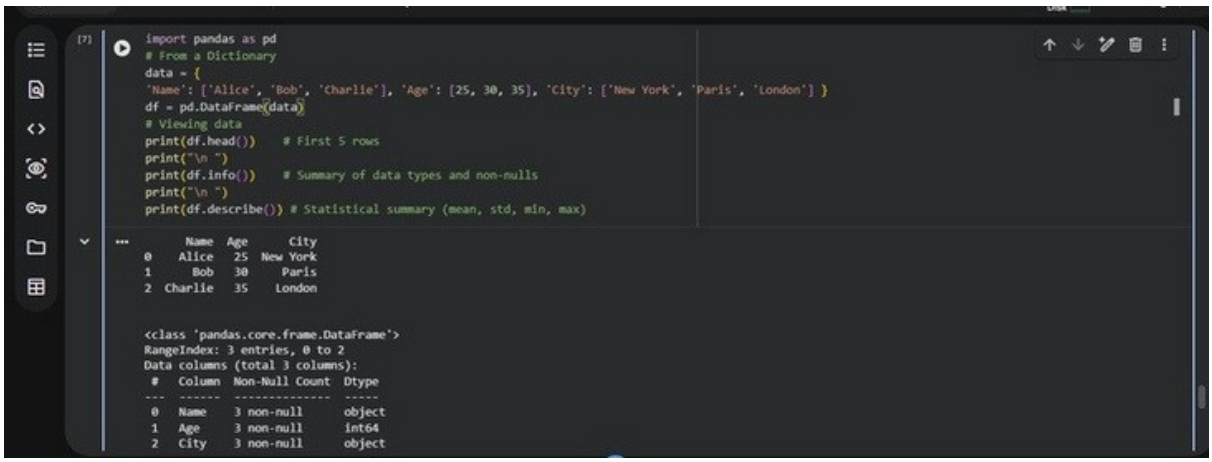
4.Convert the dictionary into a DataFrame using pd.DataFrame().

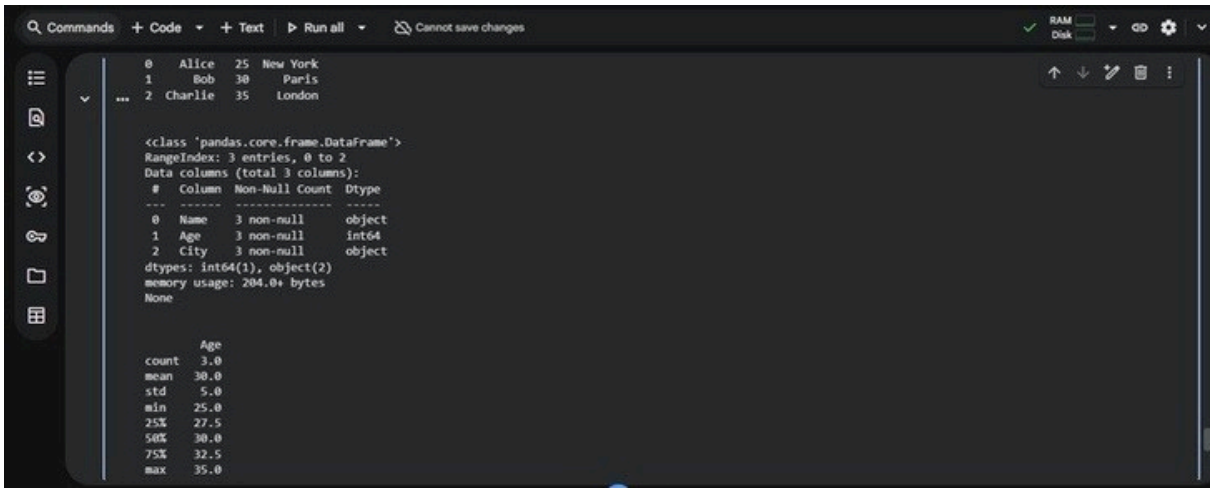
5.Display the first five rows using df.head().

6.Display the DataFrame information using df.info().

7.Display the statistical summary using df.describe().

8.Stop the program.





Result

Thus, the Python program to create a Pandas DataFrame and display its contents, data types, and statistical summary was successfully executed.

B. Selection and Filtering

Aim

To write a Python program using Pandas to select specific columns, filter rows based on conditions, and access data using iloc and loc.

Algorithm

1.Start the program.

2.Import the Pandas library as pd.

3.Use the DataFrame created in Program A.

4.Select the Age column using df['Age'].

5.Filter rows where Age is greater than 25.

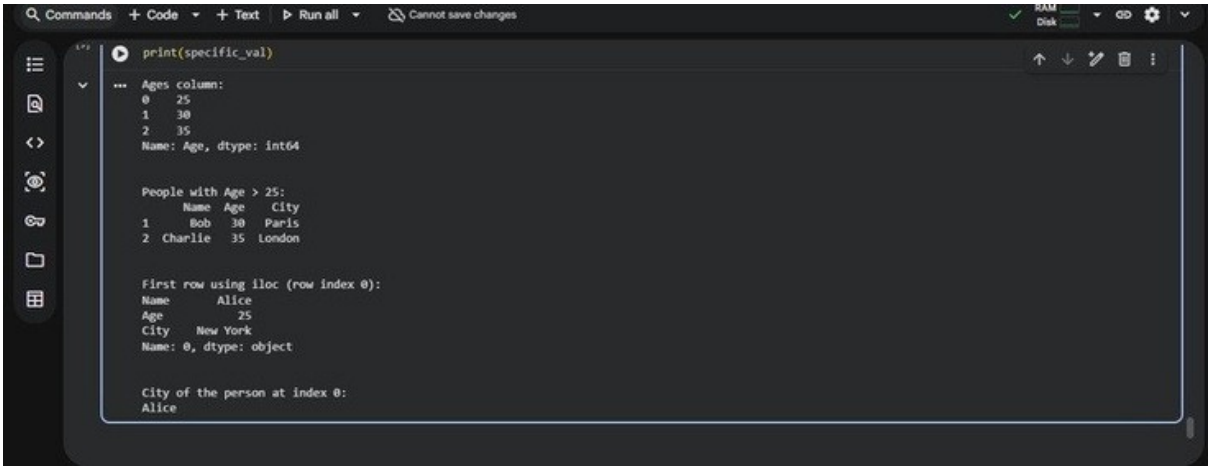
6.Access the first row using df.iloc[0].

7.Access the Name value at index 0 using df.loc[0, 'Name'].

8.Display the selected column, filtered rows, and accessed values.

9.Stop the program.





Result

Thus, the Python program to select columns, filter rows, and access DataFrame elements using iloc and loc was successfully executed.

Note: In your code, the statement df.loc[0, 'Name'] retrieves Alice, not the city. To retrieve the city, use df.loc[0, 'City'].

C. Data Cleaning (Handling Missing Values)

Aim

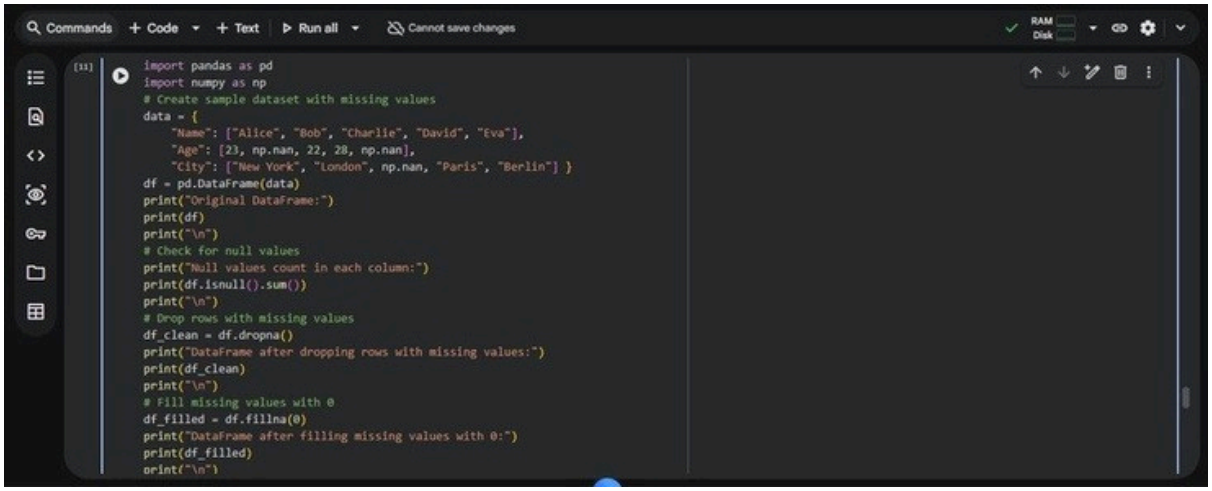
To write a Python program using Pandas and NumPy to identify and handle missing values in a DataFrame by dropping rows, filling missing values with zero, and replacing missing Age values with the column mean.

Algorithm

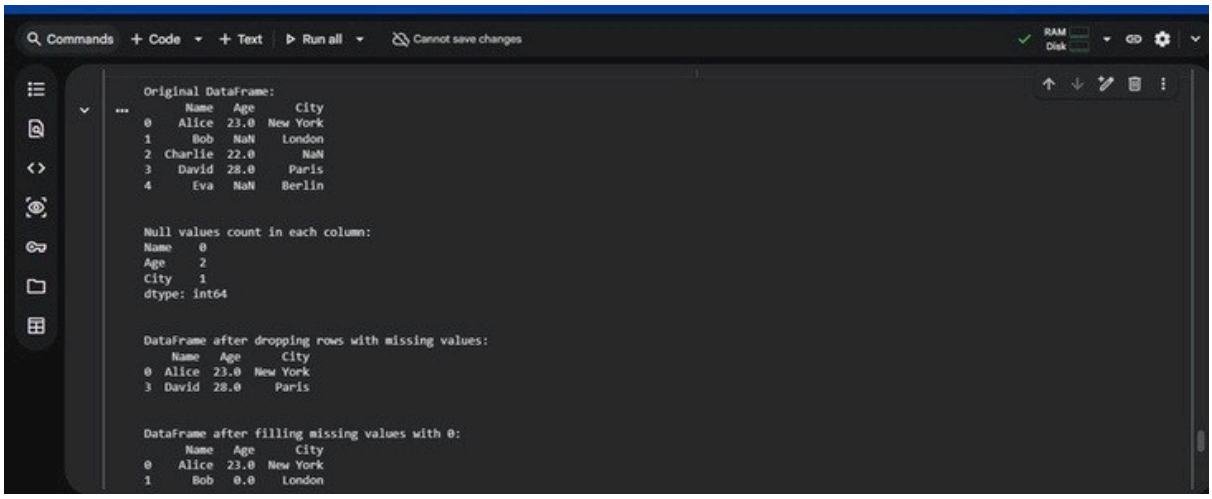
1. Start the program.
2. Import the Pandas library as pd and the NumPy library as np.
3. Create a dictionary containing Name, Age, and City, with some missing values represented by np.nan.
4. Convert the dictionary into a DataFrame using pd.DataFrame().
5. Display the original DataFrame.
6. Identify missing values in each column using df.isnull().sum().
7. Remove rows containing missing values using df.dropna().
8. Fill all missing values with zero using df.fillna(0).
9. Calculate the mean of the Age column and replace missing Age values with the mean.

10. Display the results of each cleaning operation.

11. Stop the program.



```
import pandas as pd
import numpy as np
# Create sample dataset with missing values
data = {
    "Name": ["Alice", "Bob", "Charlie", "David", "Eva"],
    "Age": [23, np.nan, 22, 28, np.nan],
    "City": ["New York", "London", np.nan, "Paris", "Berlin"] }
df = pd.DataFrame(data)
print("Original DataFrame:")
print(df)
print("\n")
# Check for null values
print("Null values count in each column:")
print(df.isnull().sum())
print("\n")
# Drop rows with missing values
df_clean = df.dropna()
print("DataFrame after dropping rows with missing values:")
print(df_clean)
print("\n")
# Fill missing values with 0
df_filled = df.fillna(0)
print("DataFrame after filling missing values with 0:")
print(df_filled)
print("\n")
```



```
Original DataFrame:
      Name   Age      City
0    Alice  23.0  New York
1      Bob   NaN    London
2  Charlie  22.0       NaN
3    David  28.0     Paris
4      Eva   NaN    Berlin


Null values count in each column:
Name    0
Age     2
City    1
dtype: int64


DataFrame after dropping rows with missing values:
    Name   Age      City
0  Alice  23.0  New York
3  David  28.0     Paris


DataFrame after filling missing values with 0:
      Name   Age      City
0    Alice  23.0  New York
1      Bob   0.0    London
```

Result

Thus, the Python program to identify and handle missing values in a Pandas DataFrame using isnull(), dropna(), and fillna() was successfully executed.

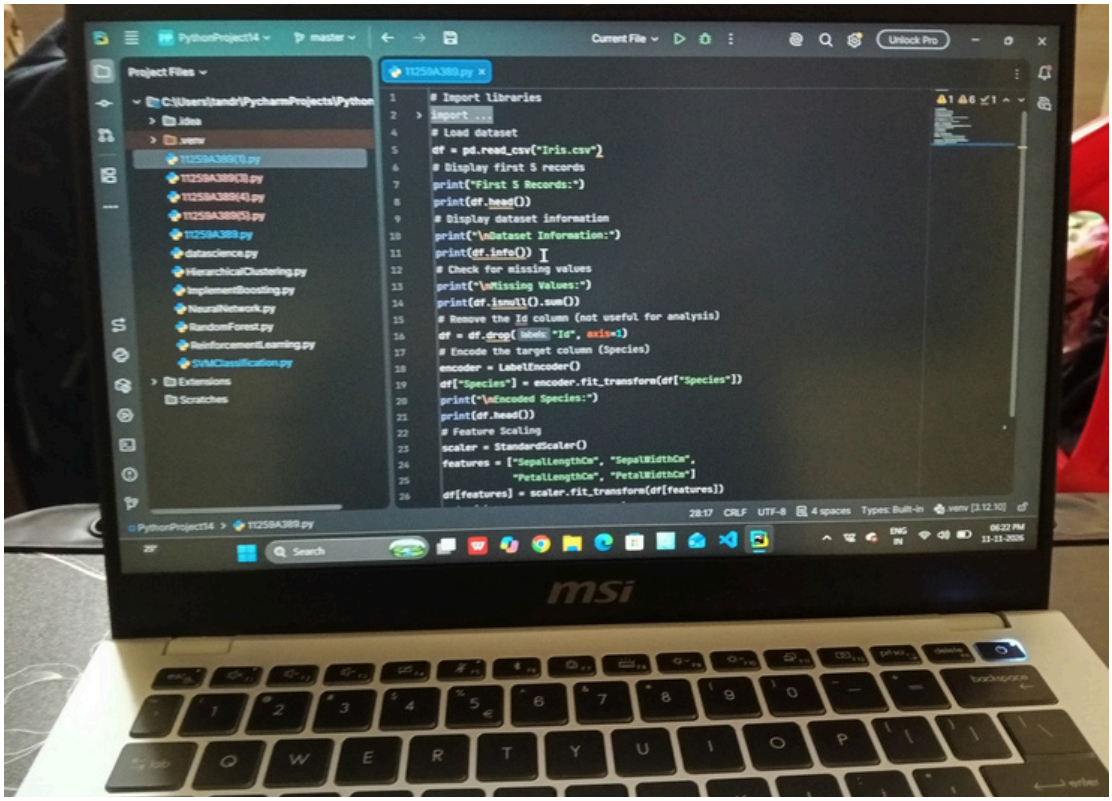
Program 1: Data Preprocessing

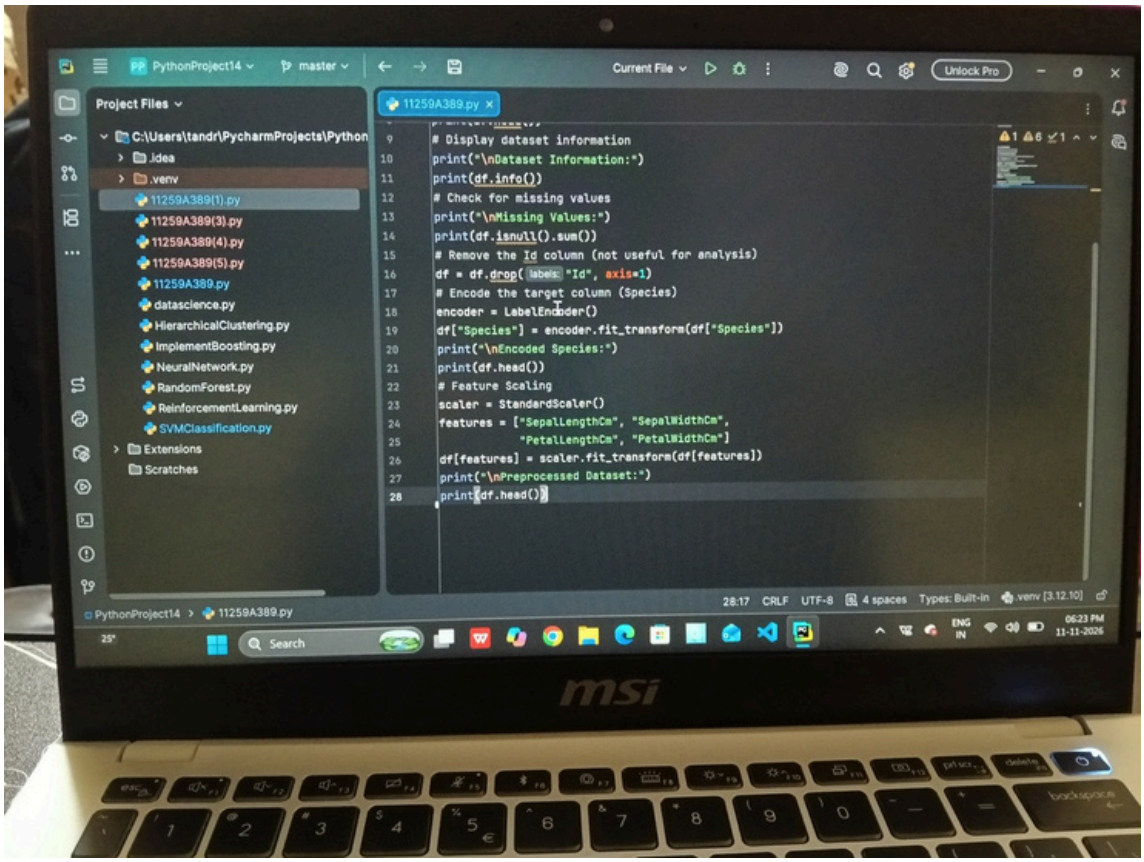
1. Aim

To preprocess a dataset by checking for missing values, filling missing numerical values with their mean, selecting numerical columns, and normalizing the data using StandardScaler.

2. Algorithm

1. Start.
2. Import the required libraries: pandas and StandardScaler from sklearn.preprocessing.
3. Load the dataset from data.csv.
4. Display the first five rows of the original dataset.
5. Check and display the number of missing values in each column.
6. Fill missing numerical values with the mean of their respective columns.
7. Select the columns Hours_Studied, Attendance, Assignments, and Marks.
8. Create a StandardScaler object.
9. Apply standardization to the selected numerical columns.
10. Convert the normalized data into a Pandas DataFrame.
11. Display the first five rows of the normalized dataset.
12. Stop.


PythonProject14
master
Current File
Unlock Pro
Project Files
11259A389.py
Import libraries
import ...
Load dataset
df = pd.read_csv("Iris.csv")
Display first 5 records
print("First 5 Records:")
print(df.head())
Display dataset information
print("\nDataset Information:")
print(df.info())
Check for missing values
print("\nMissing Values:")
print(df.isnull().sum())
Remove the Id column (not useful for analysis)
df = df.drop(labels: "Id", axis=1)
Encode the target column (Species)
encoder = LabelEncoder()
df["Species"] = encoder.fit_transform(df["Species"])
print("\nEncoded Species:")
print(df.head())
Feature Scaling
scaler = StandardScaler()
features = ["SepalLengthCm", "SepalWidthCm",
"PetalLengthCm", "PetalWidthCm"]
df[features] = scaler.fit_transfrom(df[features])
28:17 CRLF UTF-8 4 spaces Types: Built-in venv [3.12.10]
Search
msi


PythonProject14
master
Current File
Unlock Pro
Project Files
C:\Users\tandr\PycharmProjects\Python
.idea
.venv
11259A389(1).py
11259A389(3).py
11259A389(4).py
11259A389(5).py
11259A389.py
datascience.py
HierarchicalClustering.py
ImplementBoosting.py
NeuralNetwork.py
RandomForest.py
ReinforcementLearning.py
SVMClassification.py
Extensions
Scratches
11259A389.py
Display dataset information
print("\nDataset Information:")
print(df.info())
Check for missing values
print("\nMissing Values:")
print(df.isnull().sum())
Remove the Id column (not useful for analysis)
df = df.drop(labels: "Id", axis=1)
Encode the target column (Species)
encoder = LabelEncoder()
df["Species"] = encoder.fit_transform(df["Species"])
print("\nEncoded Species:")
print(df.head())
Feature Scaling
scaler = StandardScaler()
features = ["SepalLengthCm", "SepalWidthCm",
"PetalLengthCm", "PetalWidthCm"]
df[features] = scaler.fit_transform(df[features])
print("\nPreprocessed Dataset:")
print(df.head())
28:17 CRLF UTF-8 4 spaces Types: Built-in venv [3.12.10]
PythonProject14 > 11259A389.py
Search
msi

```
First 5 Records:
   Id  SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm      Species
0   1            5.1           3.5            1.4           0.2  Iris-setosa
1   2            4.9           3.0            1.4           0.2  Iris-setosa
2   3            4.7           3.2            1.3           0.2  Iris-setosa
3   4            4.6           3.1            1.5           0.2  Iris-setosa
4   5            5.0           3.6            1.4           0.2  Iris-setosa

Dataset Information:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
 #   Column         Non-Null Count  Dtype
---  ------         --------------  -----
 0   Id             150 non-null    int64
 1   SepalLengthCm  150 non-null    float64
 2   SepalWidthCm   150 non-null    float64
 3   PetalLengthCm  150 non-null    float64
 4   PetalWidthCm   150 non-null    float64
 5   Species        150 non-null    object
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
None

Missing Values:
Id               0
SepalLengthCm    0
SepalWidthCm     0
PetalLengthCm    0
PetalWidthCm     0
Species          0
```

result

Thus, the data preprocessing program was implemented successfully. It loads the dataset, identifies and handles missing numerical values, selects the required features, and standardizesthem using StandardScaler, preparing the data for further machine learning analysis.

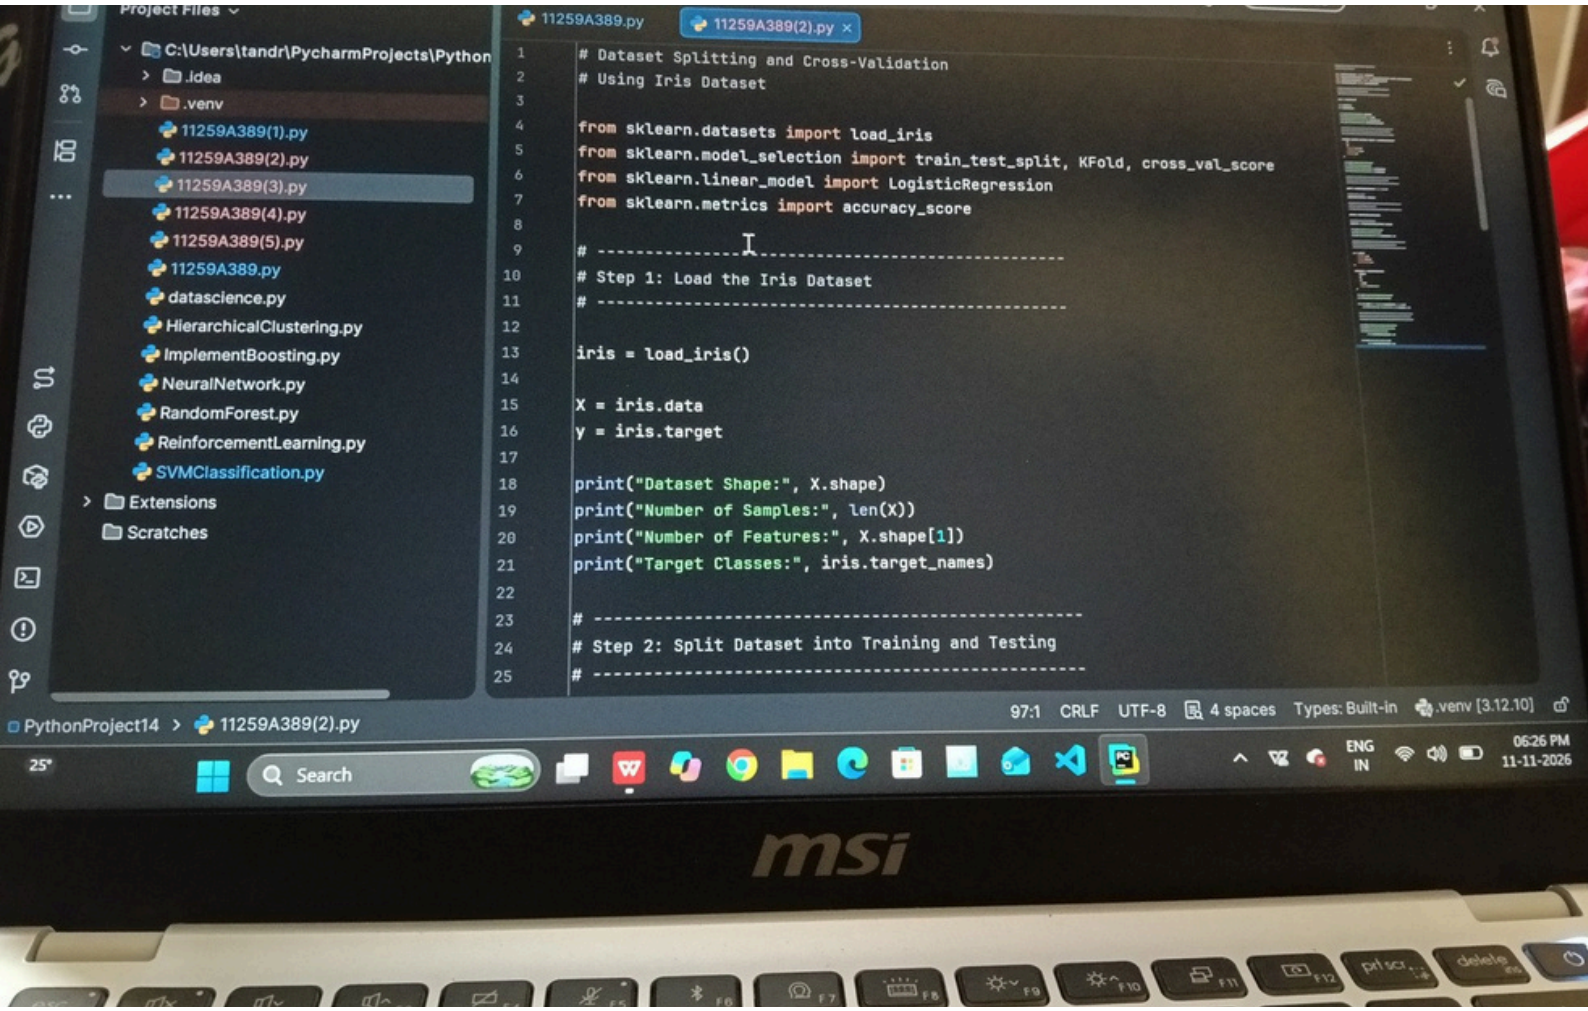
Data Split-Up and Cross-Validation

Aim

To split the dataset into training and testing sets and evaluate a K-Nearest Neighbors (KNN) classifier using 5-fold cross-validation.

Algorithm

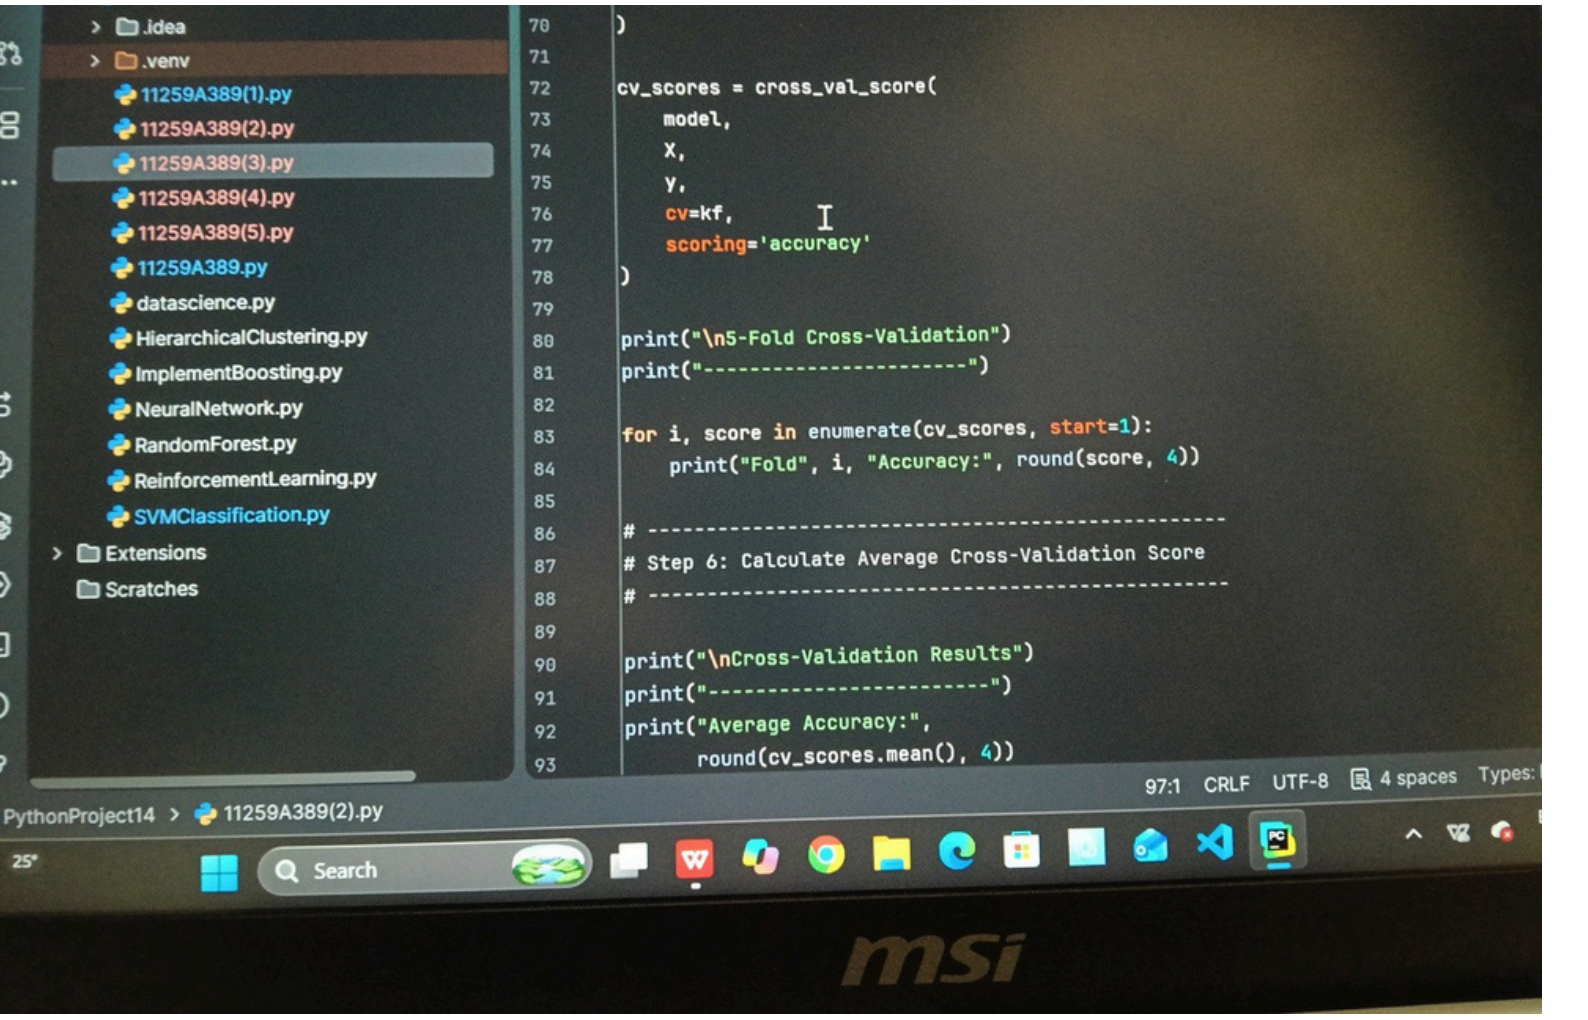
1. Start.
2. Import pandas, train_test_split, cross_val_score, and KNeighborsClassifier.
3. Load the dataset from data.csv.
4. Select the input features: Hours_Studied, Attendance, and Assignments.
5. Select Result as the target variable.
6. Split the dataset into 80% training data and 20% testing data.
7. Display the sizes of the training and testing datasets.
8. Create a KNN classifier with n_neighbors=3.
9. Perform 5-fold cross-validation using cross_val_score().
10. Display the accuracy scores obtained from the five folds.
11. Calculate and display the average accuracy.
12. Stop.


Project Files
C:\Users\tandr\PycharmProjects\Python
.idea
.venv
11259A389(1).py
11259A389(2).py
11259A389(3).py
11259A389(4).py
11259A389(5).py
11259A389.py
datascience.py
HierarchicalClustering.py
ImplementBoosting.py
NeuralNetwork.py
RandomForest.py
ReinforcementLearning.py
SVMClassification.py
Extensions
Scratches
11259A389.py
11259A389(2).py
Dataset Splitting and Cross-Validation
Using Iris Dataset
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split, KFold, cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

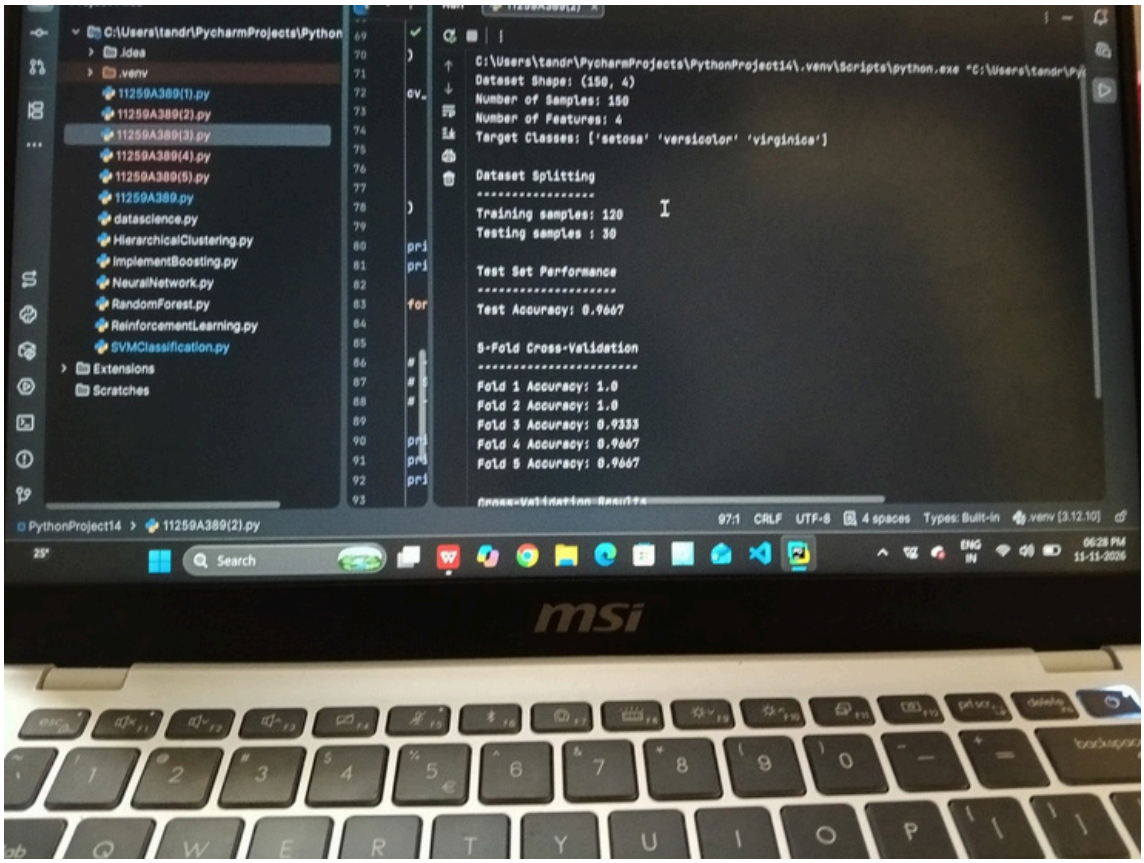
Step 1: Load the Iris Dataset
----
iris = load_iris()
X = iris.data
y = iris.target
print("Dataset Shape:", X.shape)
print("Number of Samples:", len(X))
print("Number of Features:", X.shape[1])
print("Target Classes:", iris.target_names)

Step 2: Split Dataset into Training and Testing

PythonProject14 > 11259A389(2).py
97:1 CRLF UTF-8 4 spaces Types: Built-in .venv [3.12.10]
Search
ENG IN
06:26 PM 11-11-2026
msi


.idea
.venv
11259A389(1).py
11259A389(2).py
11259A389(3).py
11259A389(4).py
11259A389(5).py
11259A389.py
datascience.py
HierarchicalClustering.py
ImplementBoosting.py
NeuralNetwork.py
RandomForest.py
ReinforcementLearning.py
SVMClassification.py
Extensions
Scratches
)
cv_scores = cross_val_score(
model,
X,
y,
cv=kf,
scoring='accuracy'
)
print("\n5-Fold Cross-Validation")
print("----------")
for i, score in enumerate(cv_scores, start=1):
print("Fold", i, "Accuracy:", round(score, 4))

Step 6: Calculate Average Cross-Validation Score
----
print("\nCross-Validation Results")
print("----------")
print("Average Accuracy:",
round(cv_scores.mean(), 4))
PythonProject14 > 11259A389(2).py
97:1 CRLF UTF-8 4 spaces Types:
Search
msi



Result:

The dataset was successfully split into training and testing sets, and the KNN classifier was evaluated using 5-fold cross-validation.

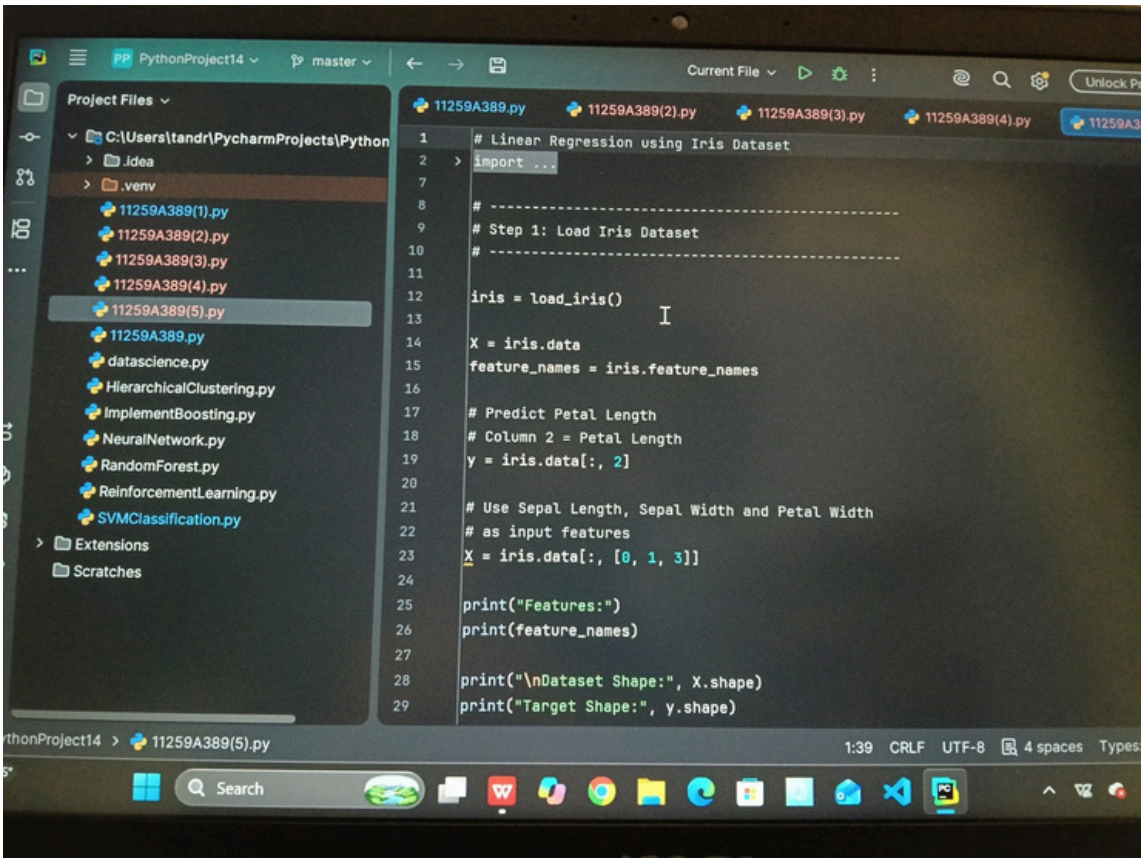
Ex.No. 3: Implement Linear Regression

Aim

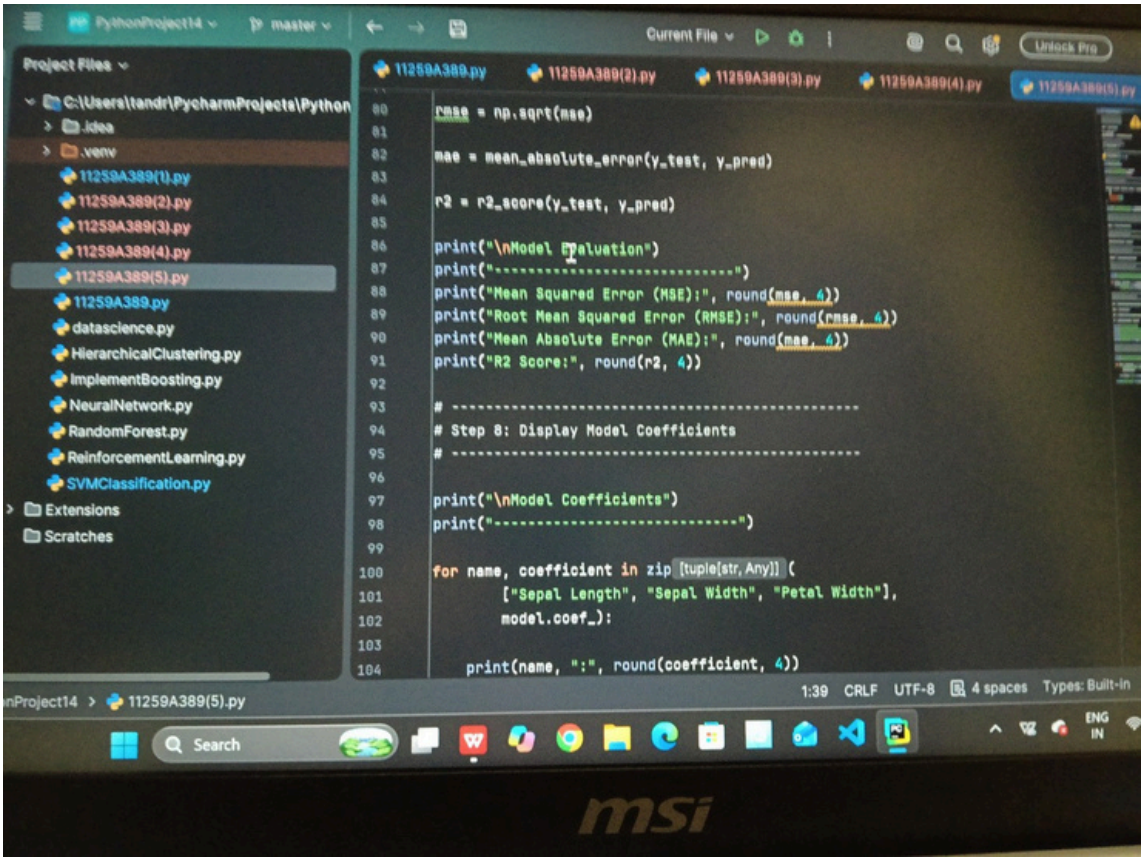
To implement a Linear Regression algorithm to predict students' marks based on the number of hours studied and evaluate the model using Mean Squared Error (MSE) and R^2 score

Algorithm

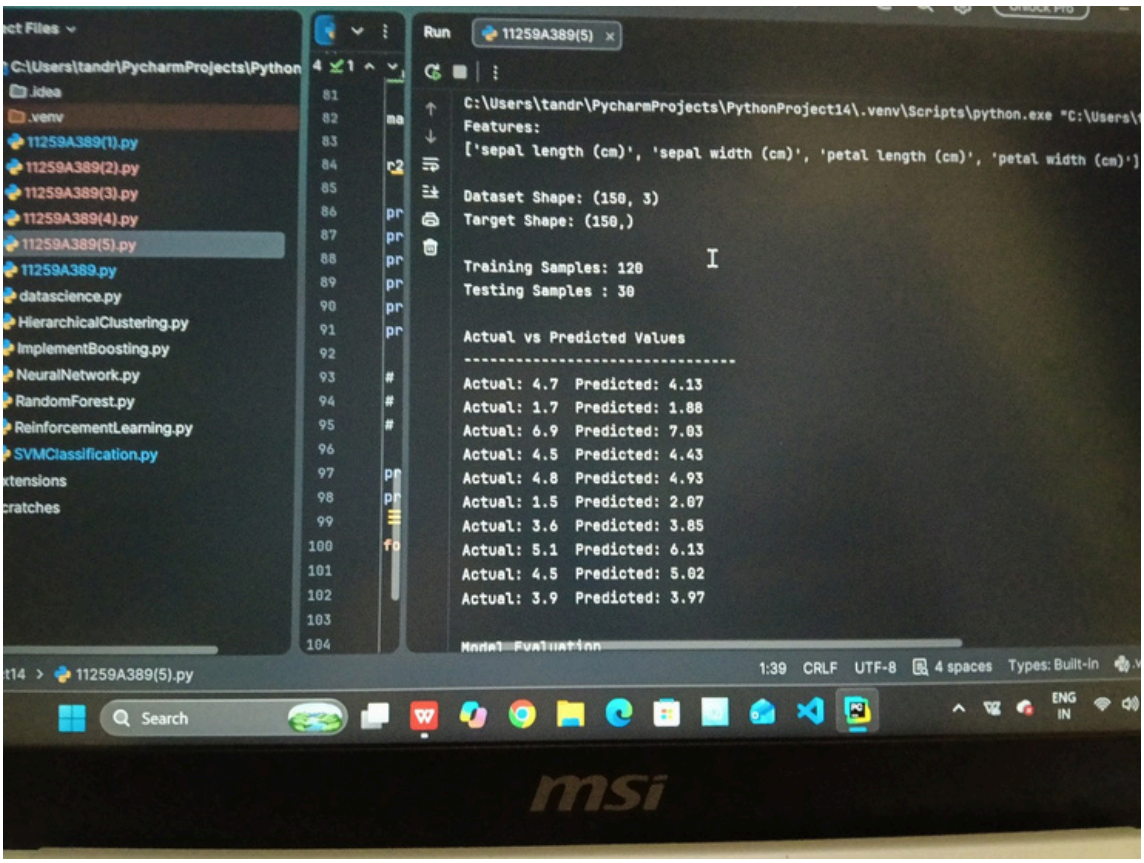
1. Start.
2. Import the required libraries: Pandas, Matplotlib, train_test_split, LinearRegression, and evaluation metrics.
3. Load the dataset from data.csv.
4. Select Hours_Studied as the input feature and Marks as the target variable.
5. Split the dataset into 80% training data and 20% testing data.
6. Create a Linear Regression model.
7. Train the model using the training dataset.
8. Predict students' marks using the testing dataset.
9. Evaluate the model using Mean Squared Error and R^2 score.
10. Display the predicted marks.
11. Plot the actual data points and the regression line.
12. Stop.

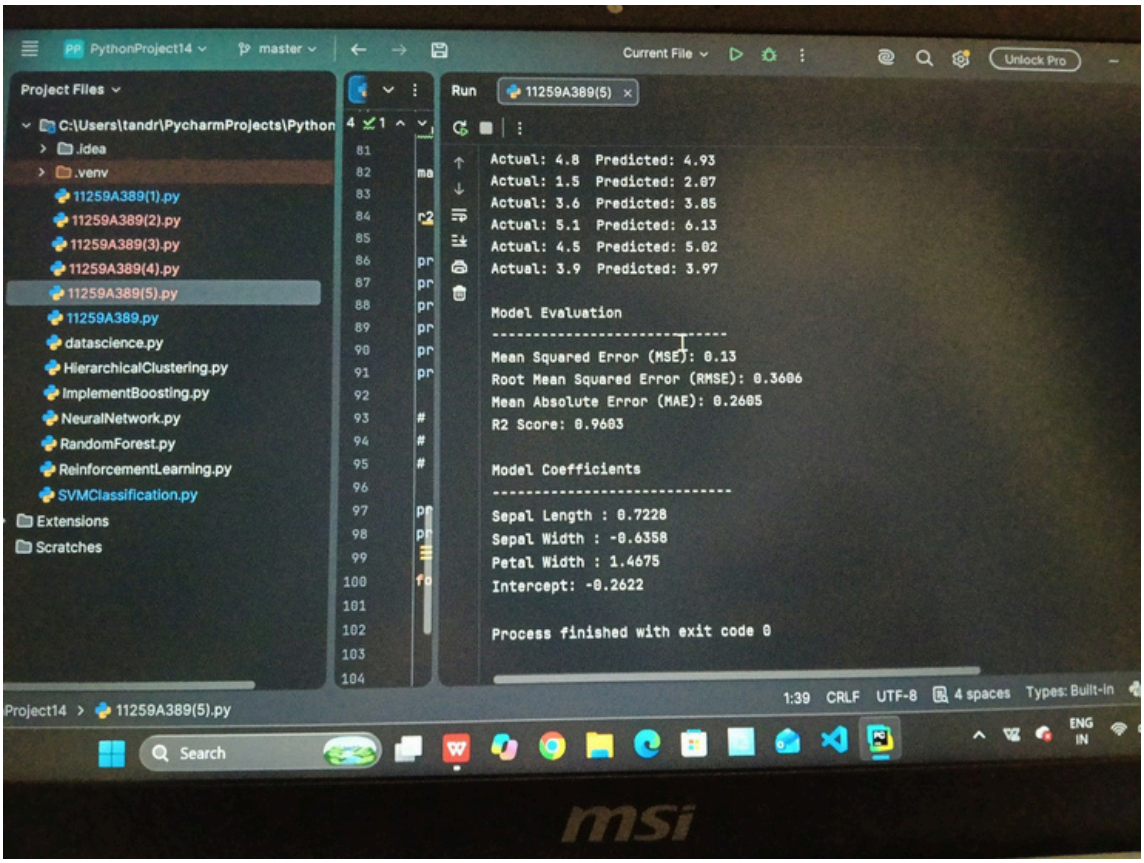

PythonProject14
master
Current File
Project Files
C:\Users\tandr\PycharmProjects\Python
.idea
.venv
11259A389(1).py
11259A389(2).py
11259A389(3).py
11259A389(4).py
11259A389(5).py
11259A389.py
datascience.py
HierarchicalClustering.py
ImplementBoosting.py
NeuralNetwork.py
RandomForest.py
ReinforcementLearning.py
SVMClassification.py
Extensions
Scratches
11259A389.py
11259A389(2).py
11259A389(3).py
11259A389(4).py
Linear Regression using Iris Dataset
import ...

Step 1: Load Iris Dataset
----
iris = load_iris()
X = iris.data
feature_names = iris.feature_names
Predict Petal Length
Column 2 = Petal Length
y = iris.data[:, 2]
Use Sepal Length, Sepal Width and Petal Width
as input features
X = iris.data[:, [0, 1, 3]]
print("Features:")
print(feature_names)
print("\nDataset Shape:", X.shape)
print("Target Shape:", y.shape)
11259A389(5).py
1:39 CRLF UTF-8 4 spaces
Search


PythonProject14
master
Current File
Project Files
C:\Users\tandr\PycharmProjects\Python
.idea
.venv
11259A389(1).py
11259A389(2).py
11259A389(3).py
11259A389(4).py
11259A389(5).py
11259A389.py
datascience.py
HierarchicalClustering.py
ImplementBoosting.py
NeuralNetwork.py
RandomForest.py
ReinforcementLearning.py
SVMClassification.py
Extensions
Scratches
11259A389.py
11259A389(2).py
11259A389(3).py
11259A389(4).py
11259A389(5).py
rmse = np.sqrt(mse)
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print("\nModel Evaluation")
print("----------")
print("Mean Squared Error (MSE):", round(mse, 4))
print("Root Mean Squared Error (RMSE):", round(rmse, 4))
print("Mean Absolute Error (MAE):", round(mae, 4))
print("R2 Score:", round(r2, 4))

Step 8: Display Model Coefficients
----
print("\nModel Coefficients")
print("----------")
for name, coefficient in zip [tuple[str, Any]] (
["Sepal Length", "Sepal Width", "Petal Width"],
model.coef_):
print(name, ":", round(coefficient, 4))
11259A389(5).py
1:39 CRLF UTF-8 4 spaces Types: Built-in
Search
msi


Run
11259A389(5)
C:\Users\tandr\PycharmProjects\PythonProject14\.venv\Scripts\python.exe "C:\Users\
Features:
['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width (cm)']
Dataset Shape: (150, 3)
Target Shape: (150,)
Training Samples: 120
Testing Samples : 30
Actual vs Predicted Values
Actual: 4.7 Predicted: 4.13
Actual: 1.7 Predicted: 1.88
Actual: 6.9 Predicted: 7.03
Actual: 4.5 Predicted: 4.43
Actual: 4.8 Predicted: 4.93
Actual: 1.5 Predicted: 2.07
Actual: 3.6 Predicted: 3.85
Actual: 5.1 Predicted: 6.13
Actual: 4.5 Predicted: 5.02
Actual: 3.9 Predicted: 3.97
Model Evaluation
11259A389(5).py
1:39 CRLF UTF-8 4 spaces Types: Built-in
Search
msi


PythonProject14
master
Current File
Unlock Pro
Project Files
Run
11259A389(5)
Actual: 4.8 Predicted: 4.93
Actual: 1.5 Predicted: 2.07
Actual: 3.6 Predicted: 3.85
Actual: 5.1 Predicted: 6.13
Actual: 4.5 Predicted: 5.02
Actual: 3.9 Predicted: 3.97
Model Evaluation
Mean Squared Error (MSE): 0.13
Root Mean Squared Error (RMSE): 0.3606
Mean Absolute Error (MAE): 0.2605
R2 Score: 0.9603
Model Coefficients
Sepal Length : 0.7228
Sepal Width : -0.6358
Petal Width : 1.4675
Intercept: -0.2622
Process finished with exit code 0
1:39 CRLF UTF-8 4 spaces Types: Built-in
Search
msi

Result

Thus, the Linear Regression algorithm was implemented successfully to predict students' marks based on their study hours. The model's performance was evaluated using Mean Squared Error and R^2 score, and the actual data and regression line were plotted.

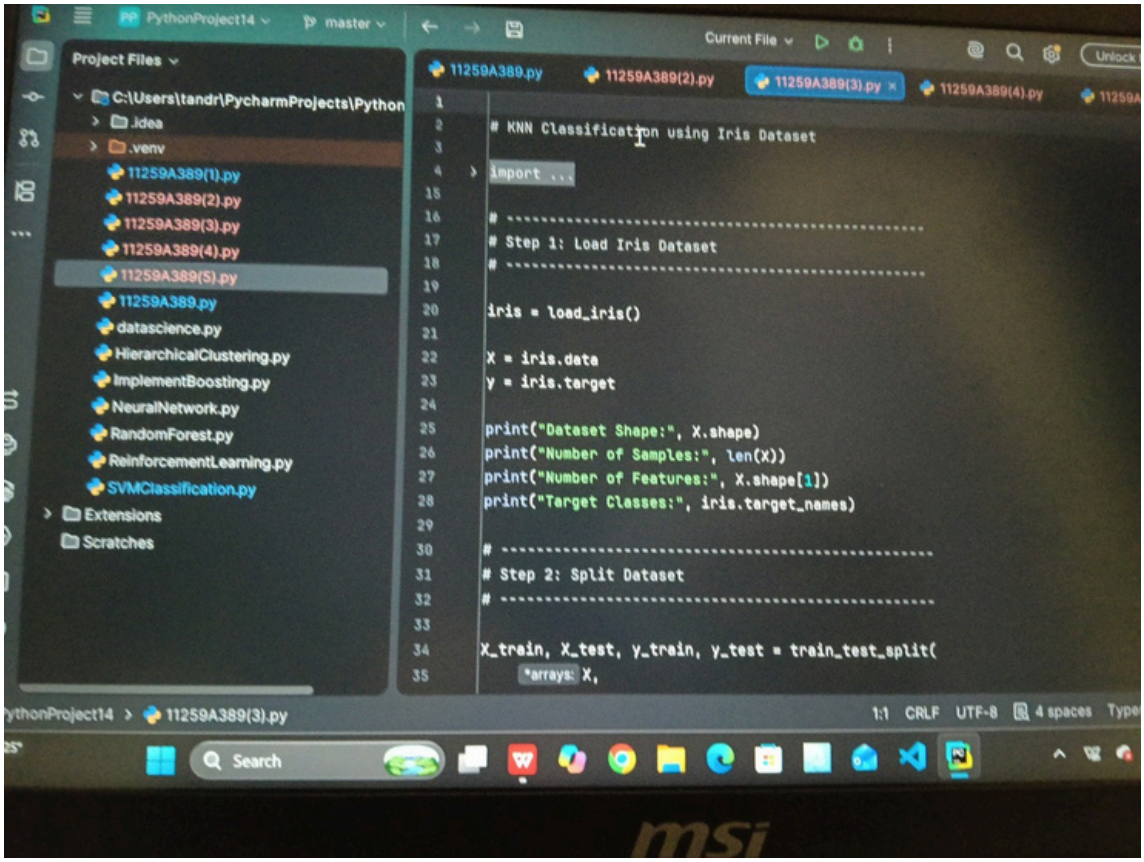
EX.NO. 4(A): Implement Classification Model Using KNN and Compute Evaluation Metrics

Aim

To implement the K-Nearest Neighbors (KNN) classification algorithm to predict students' results based on hours studied, attendance, and assignments, and evaluate the model using accuracy, confusion matrix, and classification report.

Algorithm

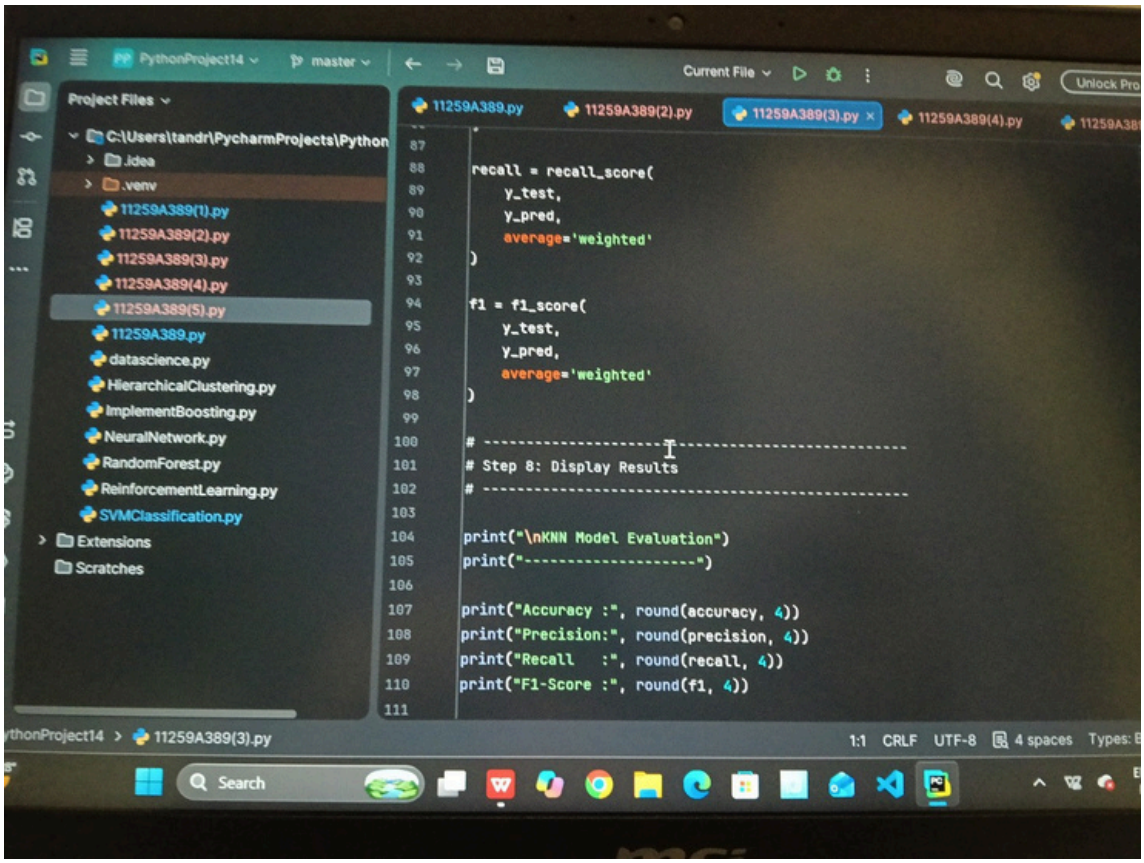
1. Start.
2. Import the required libraries: Pandas, train_test_split, KNeighborsClassifier, and evaluation metrics.
3. Load the dataset from data.csv.
4. Select Hours_Studied, Attendance, and Assignments as input features.
5. Select Result as the target variable.
6. Split the dataset into 80% training data and 20% testing data.
7. Create a KNN classifier with n_neighbors=3.
8. Train the model using the training dataset.
9. Predict the results for the testing dataset.
10. Calculate and display the accuracy score.
11. Display the confusion matrix.
12. Display the classification report containing precision, recall, and F1-score.
13. Stop.


PythonProject14
master
Current File
Project Files
C:\Users\tandr\PycharmProjects\Python
.idea
.venv
11259A389(1).py
11259A389(2).py
11259A389(3).py
11259A389(4).py
11259A389(5).py
11259A389.py
datascience.py
HierarchicalClustering.py
ImplementBoosting.py
NeuralNetwork.py
RandomForest.py
ReinforcementLearning.py
SVMClassification.py
Extensions
Scratches
11259A389.py
11259A389(2).py
11259A389(3).py
11259A389(4).py
KNN Classification using Iris Dataset
import ...

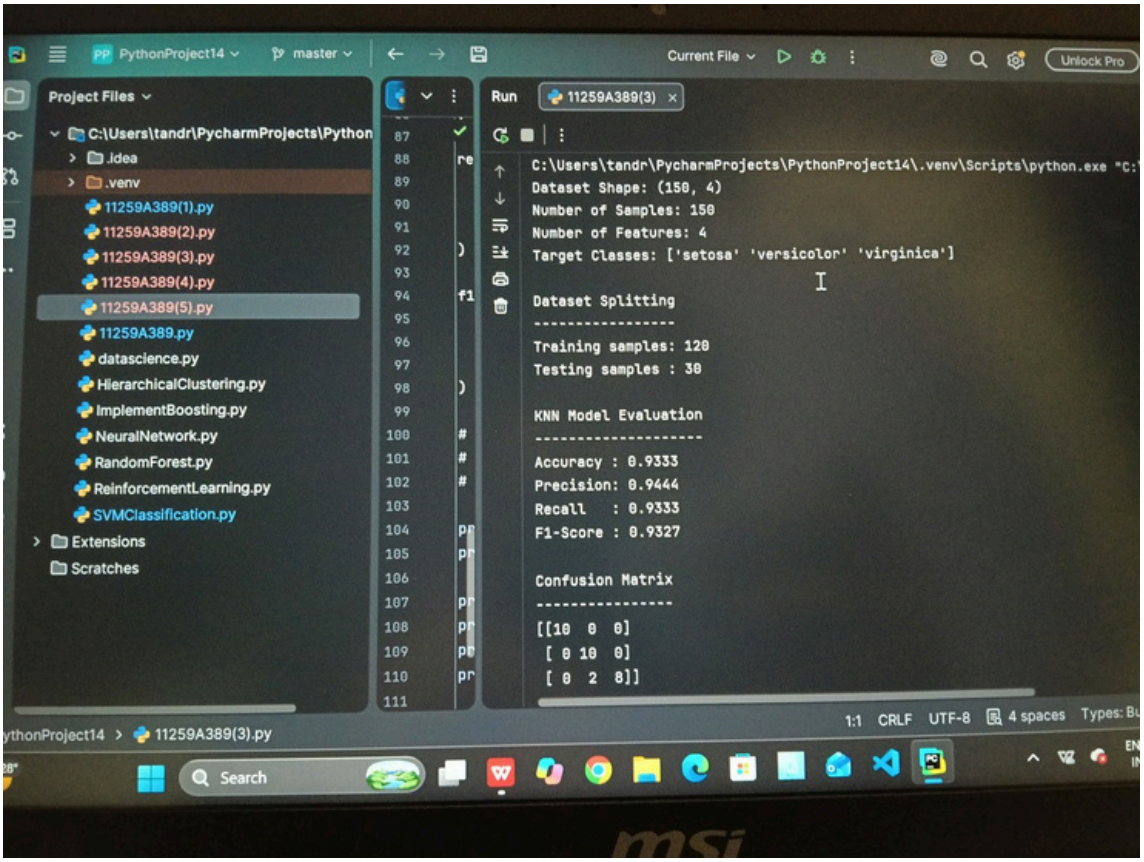
Step 1: Load Iris Dataset
----
iris = load_iris()
X = iris.data
y = iris.target
print("Dataset Shape:", X.shape)
print("Number of Samples:", len(X))
print("Number of Features:", X.shape[1])
print("Target Classes:", iris.target_names)

Step 2: Split Dataset

X_train, X_test, y_train, y_test = train_test_split(
*arrays: X,
11259A389(3).py
1:1 CRLF UTF-8 4 spaces
Search


PythonProject14
master
Current File
Project Files
C:\Users\tandr\PycharmProjects\Python
.idea
.venv
11259A389(1).py
11259A389(2).py
11259A389(3).py
11259A389(4).py
11259A389(5).py
11259A389.py
datascience.py
HierarchicalClustering.py
ImplementBoosting.py
NeuralNetwork.py
RandomForest.py
ReinforcementLearning.py
SVMClassification.py
Extensions
Scratches
11259A389.py
11259A389(2).py
11259A389(3).py
11259A389(4).py
recall = recall_score(
y_test,
y_pred,
average='weighted'
)
f1 = f1_score(
y_test,
y_pred,
average='weighted'
)

Step 8: Display Results
----
print("\nKNN Model Evaluation")
print("--------------------")
print("Accuracy :", round(accuracy, 4))
print("Precision:", round(precision, 4))
print("Recall :", round(recall, 4))
print("F1-Score :", round(f1, 4))
11259A389(3).py
1:1 CRLF UTF-8 4 spaces
Search



Result

Thus, the KNN classification model was implemented successfully, and its performance was evaluated using accuracy, confusion matrix, precision, recall, and F1-score.

EX.NO. 4(B): Implement Classification Model Using SVM and Compute Evaluation Metrics

Aim

To implement a Support Vector Machine (SVM) classification algorithm using a linear kernel to predict students' results based on hours studied, attendance, and assignments, and evaluate its performance using accuracy, confusion matrix, and classification report.

Algorithm

1. Start.
2. Import the required libraries: Pandas, train_test_split, SVC, and evaluation metrics.
3. Load the dataset from data.csv.
4. Select Hours_Studied, Attendance, and Assignments as input features.
5. Select Result as the target variable.

6.Split the dataset into 80% training data and 20% testing data.

7.Create an SVM classifier using SVC(kernel='linear').

8.Train the model using the training dataset.

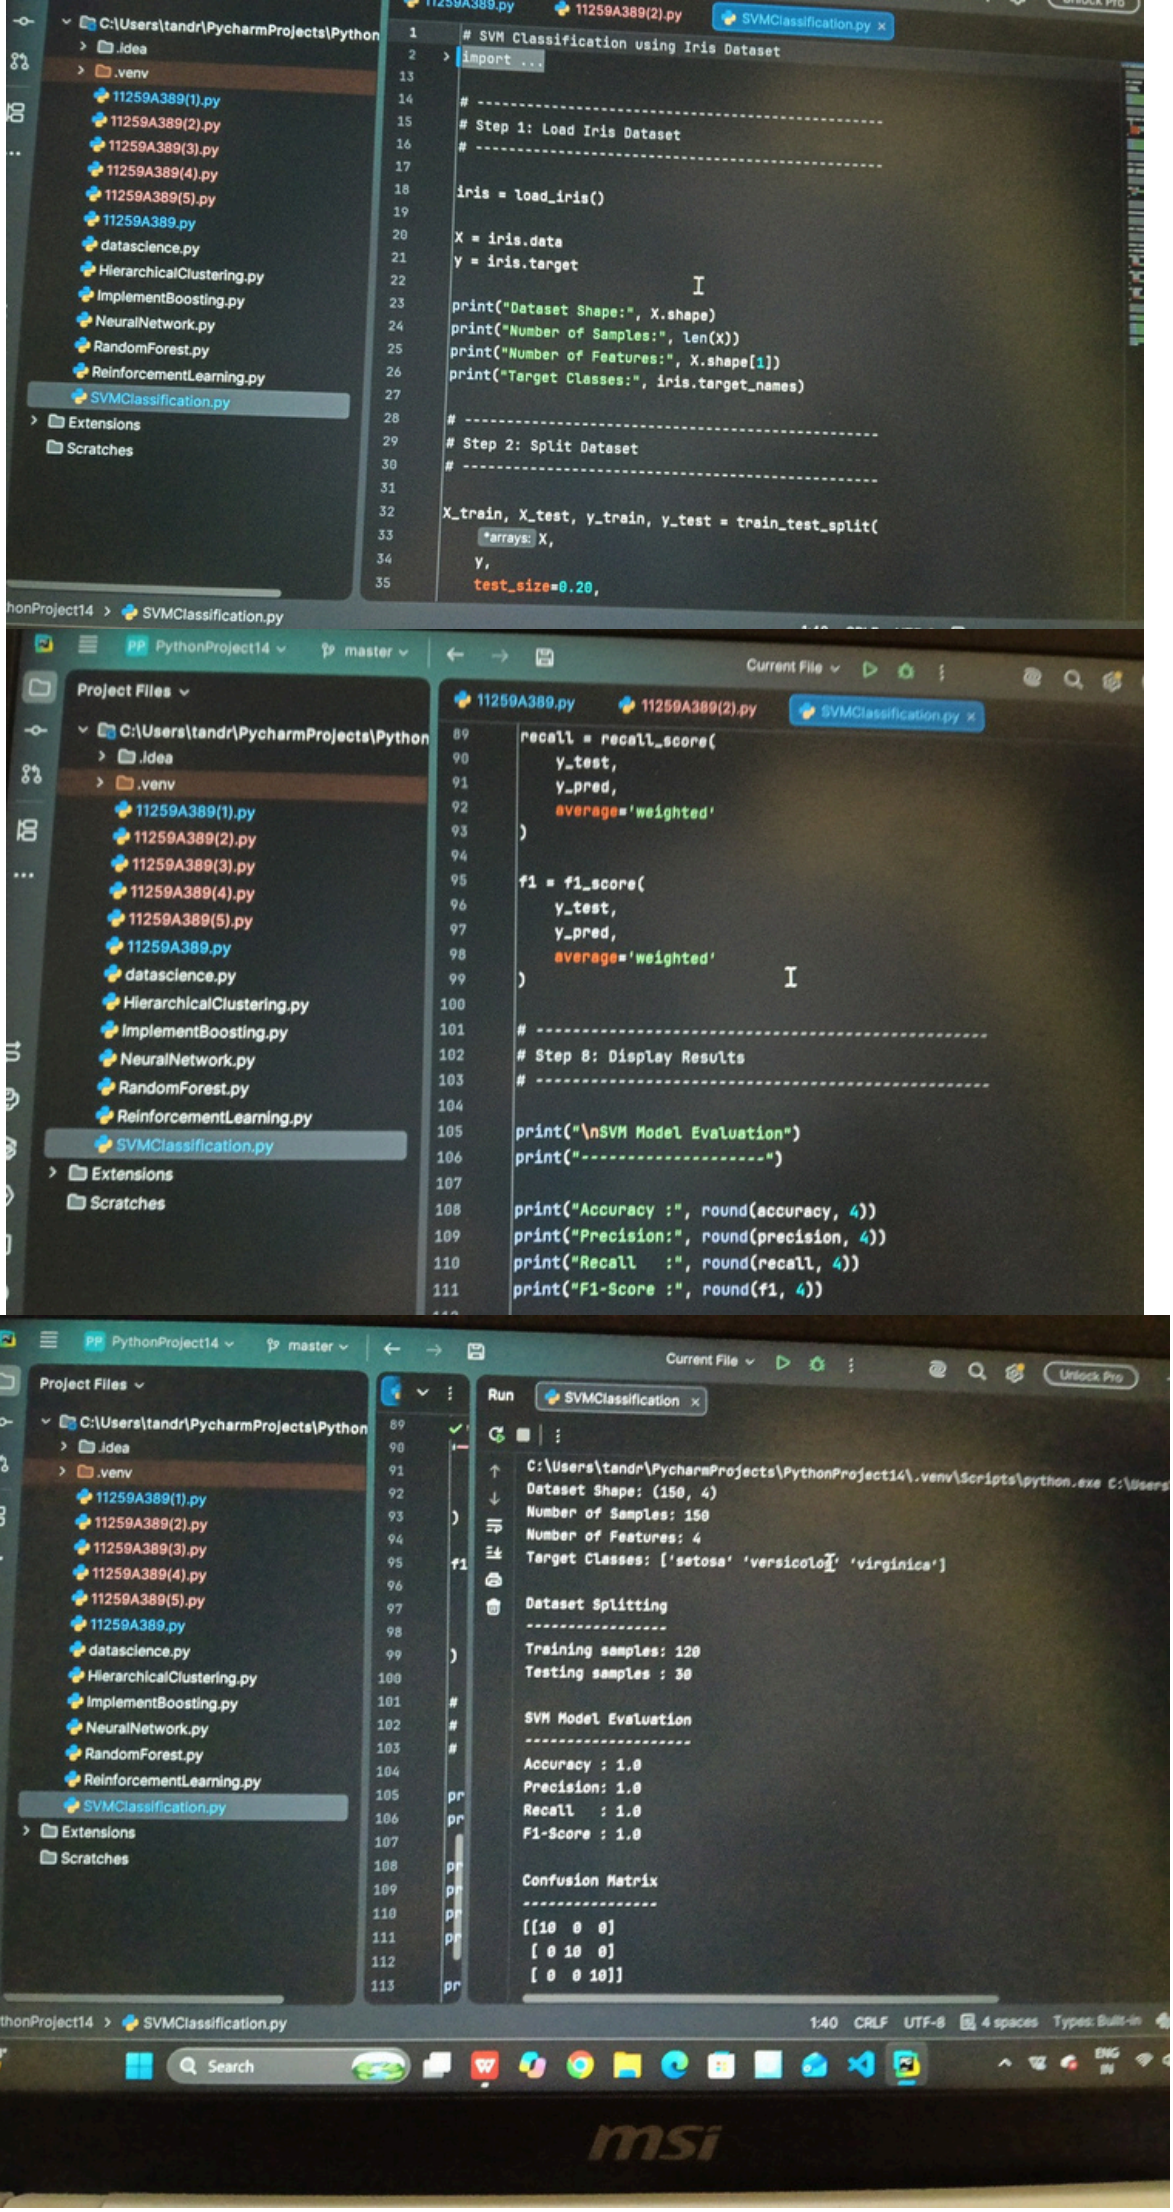
9.Predict the results for the testing dataset.

10. Calculate and display the accuracy score.

11. Display the confusion matrix.

12. Display the classification report containing precision, recall, and F1-score.

13. Stop.


C:\Users\tandr\PycharmProjects\Python
.idea
.venv
11259A389(1).py
11259A389(2).py
11259A389(3).py
11259A389(4).py
11259A389(5).py
11259A389.py
datascience.py
HierarchicalClustering.py
ImplementBoosting.py
NeuralNetwork.py
RandomForest.py
ReinforcementLearning.py
SVMClassification.py
Extensions
Scratches
SVM Classification using Iris Dataset
import ...
Step 1: Load Iris Dataset
iris = load_iris()
X = iris.data
y = iris.target
print("Dataset Shape:", X.shape)
print("Number of Samples:", len(X))
print("Number of Features:", X.shape[1])
print("Target Classes:", iris.target_names)
Step 2: Split Dataset
X_train, X_test, y_train, y_test = train_test_split(
*arrays: X,
y,
test_size=0.20,
PythonProject14
master
Current File
Project Files
recall = recall_score(
y_test,
y_pred,
average='weighted'
)
f1 = f1_score(
y_test,
y_pred,
average='weighted'
)
Step 8: Display Results
print("\nSVM Model Evaluation")
print("--------------------")
print("Accuracy :", round(accuracy, 4))
print("Precision:", round(precision, 4))
print("Recall :", round(recall, 4))
print("F1-Score :", round(f1, 4))
Run
SVMClassification
C:\Users\tandr\PycharmProjects\PythonProject14\.venv\Scripts\python.exe C:\Users
Dataset Shape: (150, 4)
Number of Samples: 150
Number of Features: 4
Target Classes: ['setosa' 'versicolor' 'virginica']
Dataset Splitting
Training samples: 120
Testing samples : 30
SVM Model Evaluation
Accuracy : 1.0
Precision: 1.0
Recall : 1.0
F1-Score : 1.0
Confusion Matrix
[[10 0 0]
[0 10 0]
[0 0 10]]
1:40
CRLF
UTF-8
4 spaces
Search
msi

Result

Thus, the SVM classification model with a linear kernel was implemented successfully, and its performance was evaluated using accuracy, confusion matrix, precision, recall, and F1-score.

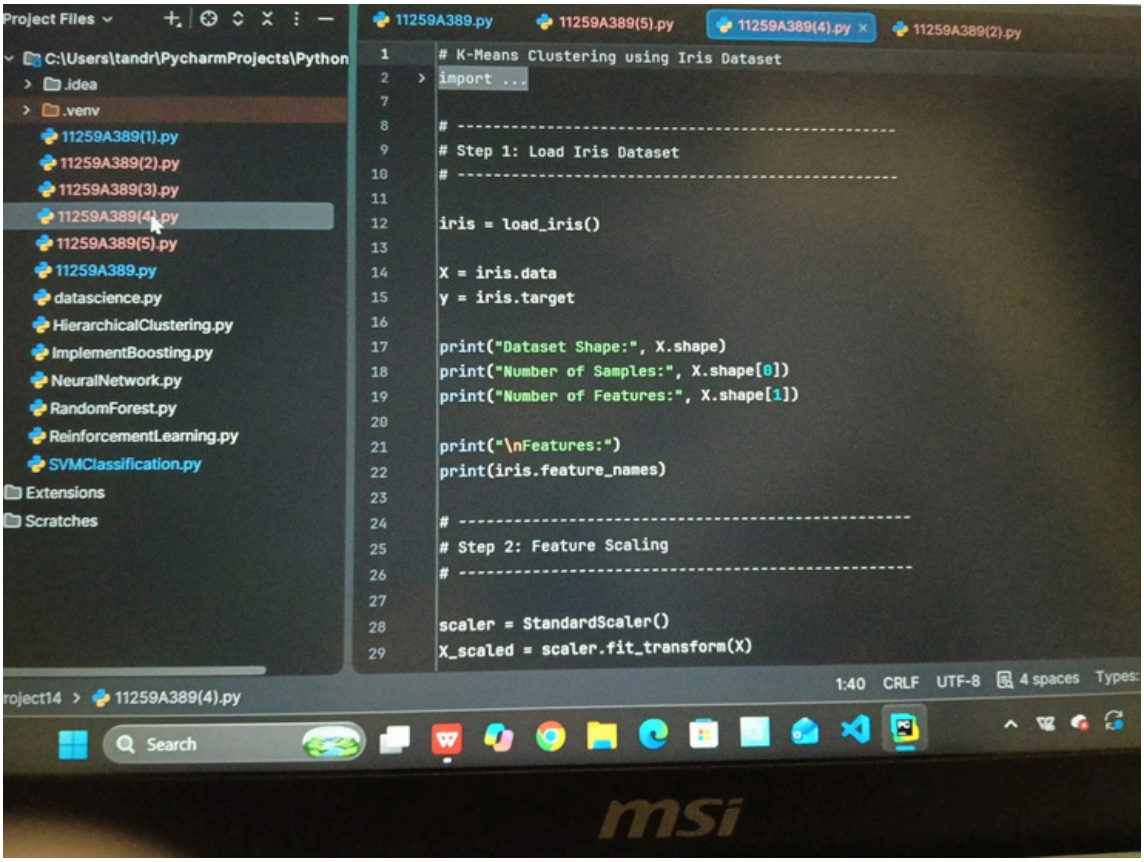
Ex No. 5: Implement K-Means Clustering

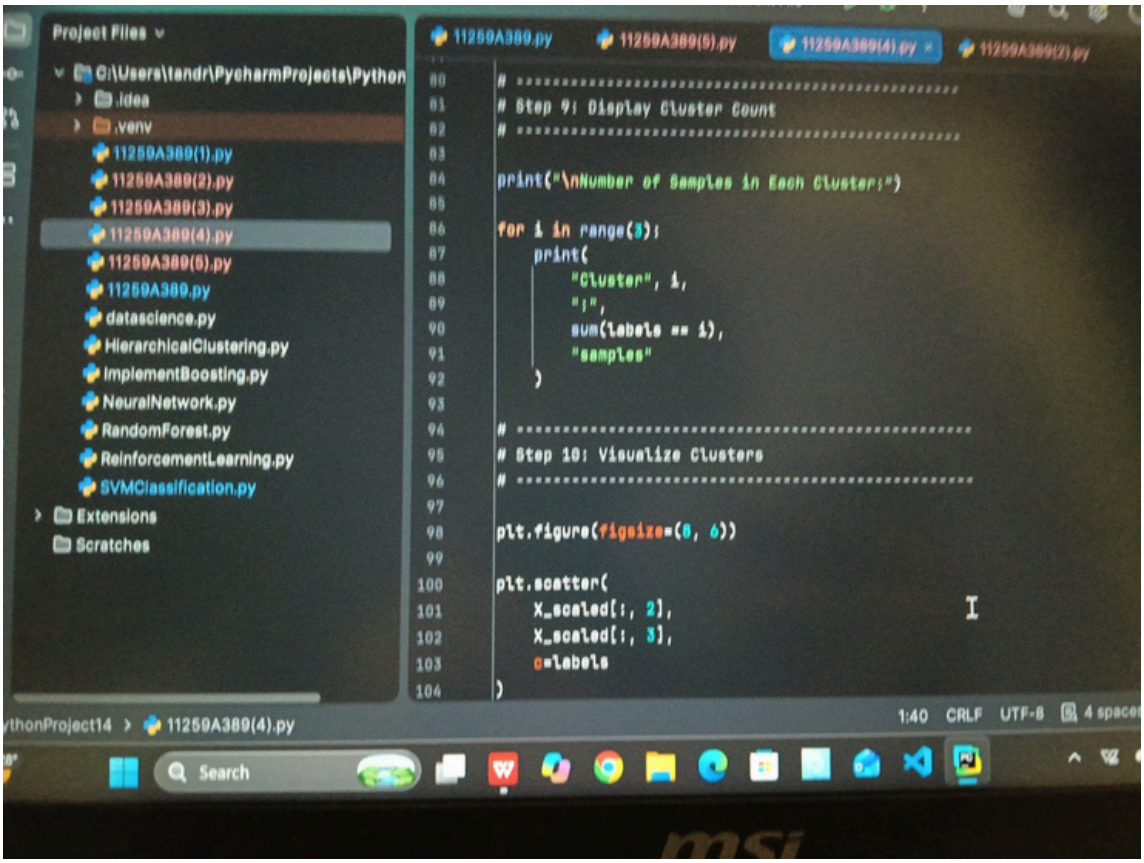
Aim

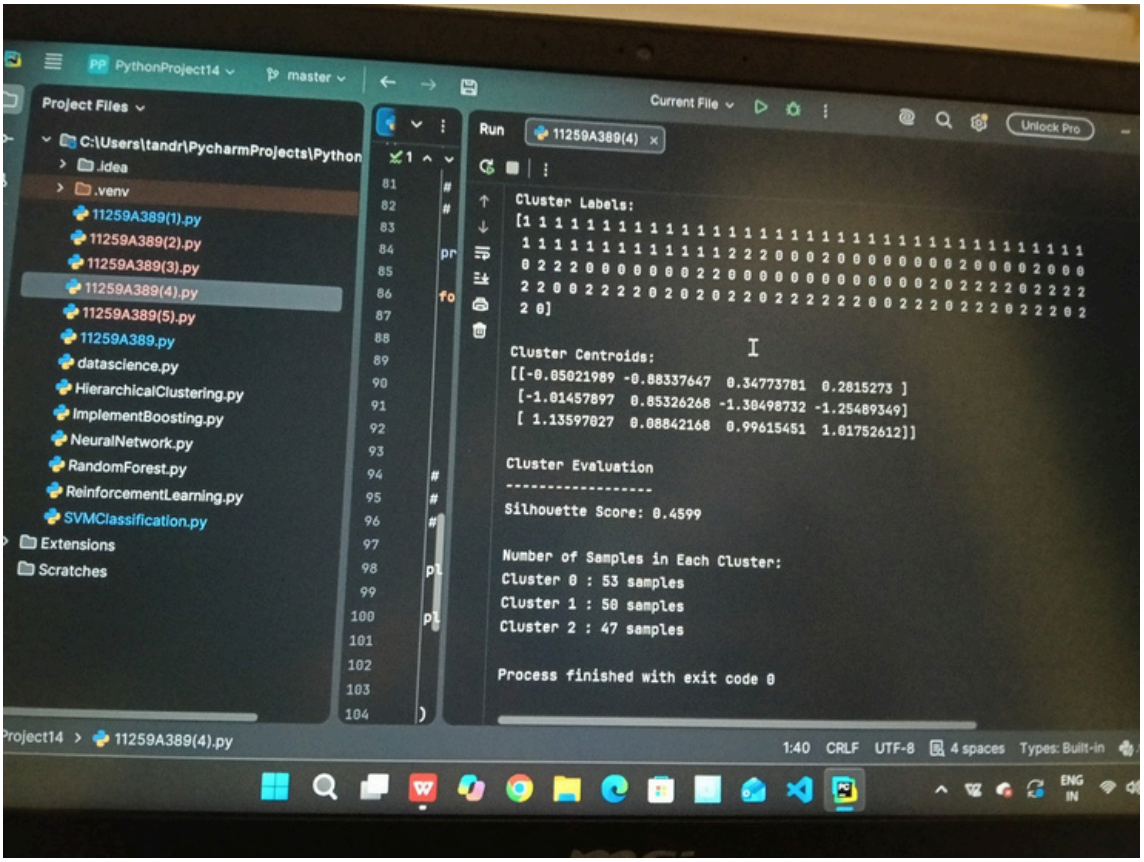
To implement the K-Means clustering algorithm to group students based on their hours studied and marks into three clusters and visualize the clusters along with their centroids.

Algorithm

1. Start.
2. Import the required libraries: Pandas, Matplotlib, and KMeans from sklearn.cluster.
3. Load the dataset from data.csv.
4. Select Hours_Studied and Marks as input features.
5. Create a K-Means clustering model with three clusters (n_clusters=3).
6. Fit the model to the selected data.
7. Obtain the cluster labels for each data point.
8. Add the cluster labels as a new column named Cluster in the dataset.
9. Display the clustered dataset.
10. Plot the data points using different colors for each cluster.
11. Plot the cluster centroids using an X-shaped marker.
12. Add axis labels, a title, and a legend to the graph.
13. Display the graph.
14. Stop


Project Files
C:\Users\tandr\PycharmProjects\Python
.idea
.venv
11259A389(1).py
11259A389(2).py
11259A389(3).py
11259A389(4).py
11259A389(5).py
11259A389.py
datascience.py
HierarchicalClustering.py
ImplementBoosting.py
NeuralNetwork.py
RandomForest.py
ReinforcementLearning.py
SVMClassification.py
Extensions
Scratches
11259A389.py
11259A389(5).py
11259A389(4).py
11259A389(2).py
K-Means Clustering using Iris Dataset
import ...
Step 1: Load Iris Dataset
iris = load_iris()
X = iris.data
y = iris.target
print("Dataset Shape:", X.shape)
print("Number of Samples:", X.shape[0])
print("Number of Features:", X.shape[1])
print("\nFeatures:")
print(iris.feature_names)
Step 2: Feature Scaling
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
1:40 CRLF UTF-8 4 spaces
11259A389(4).py
Search
msi


Project Files
C:\Users\tandr\PycharmProjects\Python
.idea
.venv
11259A389(1).py
11259A389(2).py
11259A389(3).py
11259A389(4).py
11259A389(5).py
11259A389.py
datascience.py
HierarchicalClustering.py
ImplementBoosting.py
NeuralNetwork.py
RandomForest.py
ReinforcementLearning.py
SVMClassification.py
Extensions
Scratches
11259A389.py
11259A389(5).py
11259A389(4).py
11259A389(2).py
Step 9: Display Cluster Count
print("\nNumber of Samples in Each Cluster:")
for i in range(3):
print(
"Cluster", i,
":",
sum(labels == i),
"samples"
)
Step 10: Visualize Clusters
plt.figure(figsize=(8, 6))
plt.scatter(
X_scaled[:, 2],
X_scaled[:, 3],
c=labels
)
1:40 CRLF UTF-8
11259A389(4).py
Search
msi


PythonProject14
master
Project Files
C:\Users\tandr\PycharmProjects\Python
.idea
.venv
11259A389(1).py
11259A389(2).py
11259A389(3).py
11259A389(4).py
11259A389(5).py
11259A389.py
datascience.py
HierarchicalClustering.py
ImplementBoosting.py
NeuralNetwork.py
RandomForest.py
ReinforcementLearning.py
SVMClassification.py
Extensions
Scratches
Run
11259A389(4)
Cluster Labels:
[1 1 1

Result

Thus, the K-Means clustering algorithm was implemented successfully to group students into three clusters based on their hours studied and marks. The clustered data and centroids were visualized using a scatter plot.

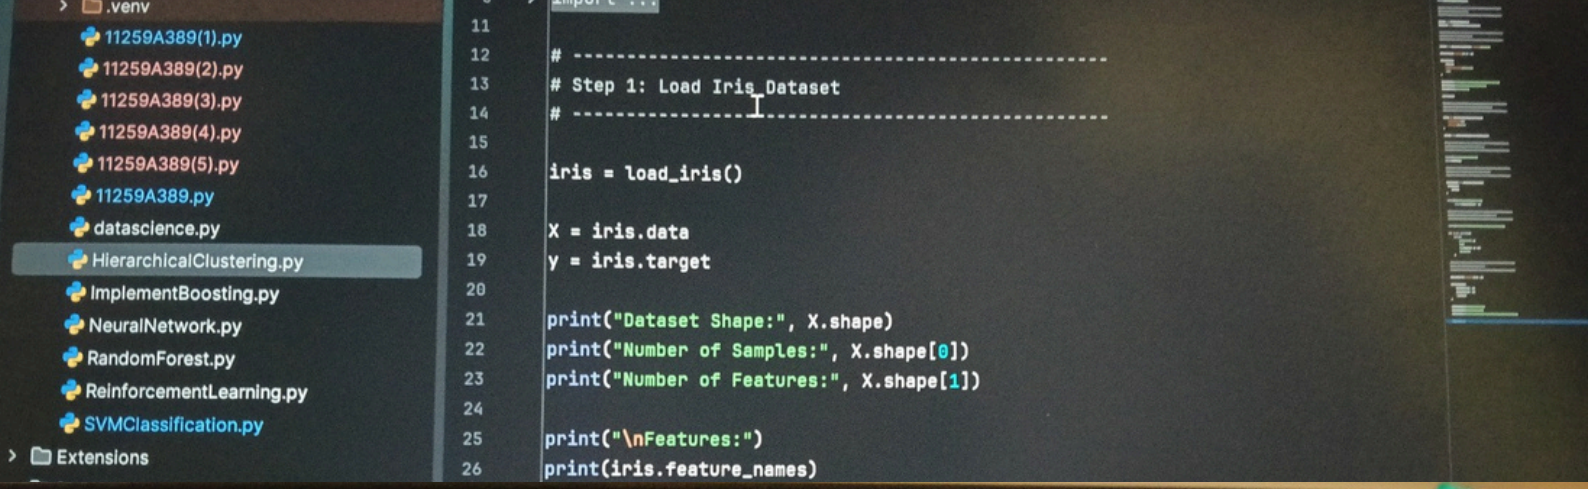
Experiment 6: Implement Hierarchical Clustering

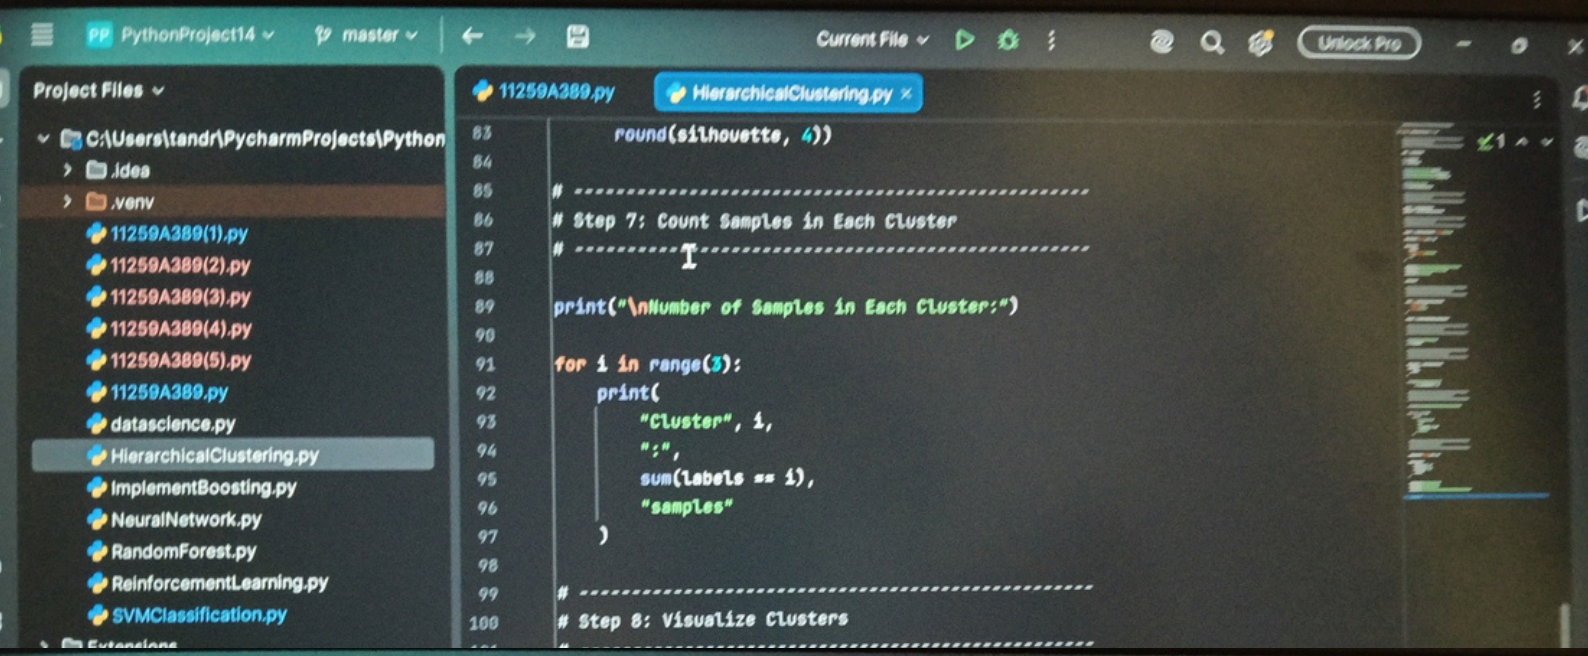
Aim

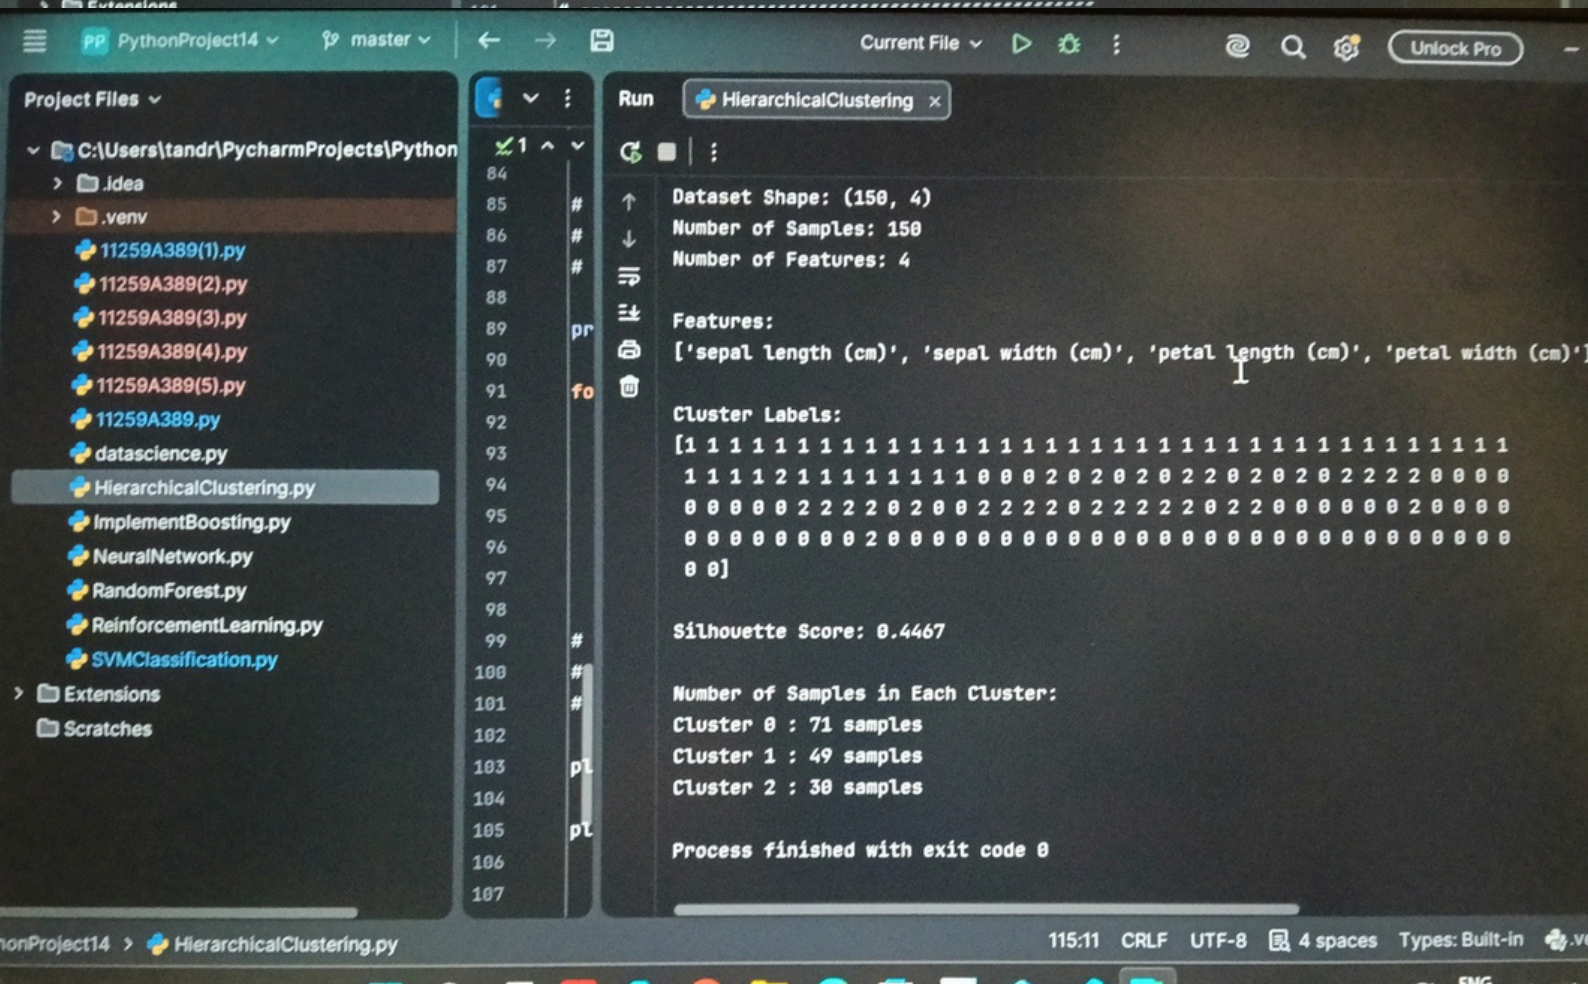
To implement the Hierarchical Clustering algorithm to group students based on their hours studied and marks into three clusters, and visualize the results using a dendrogram and a scatter plot.

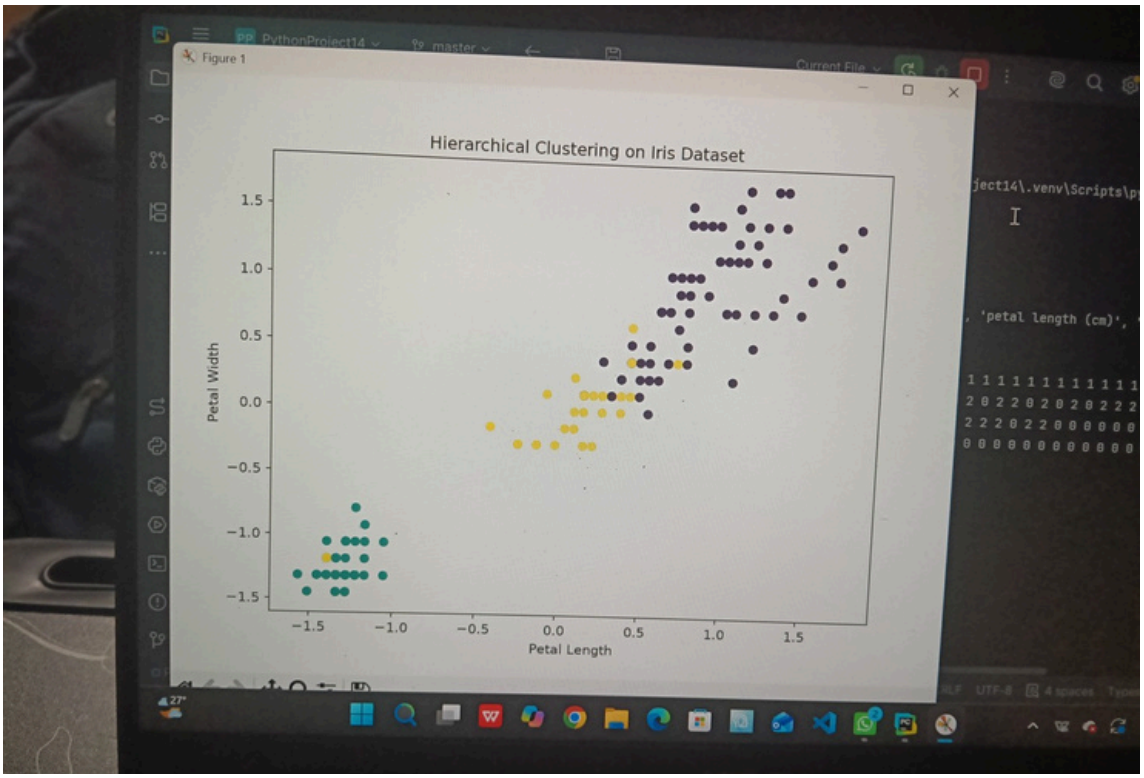
Algorithm

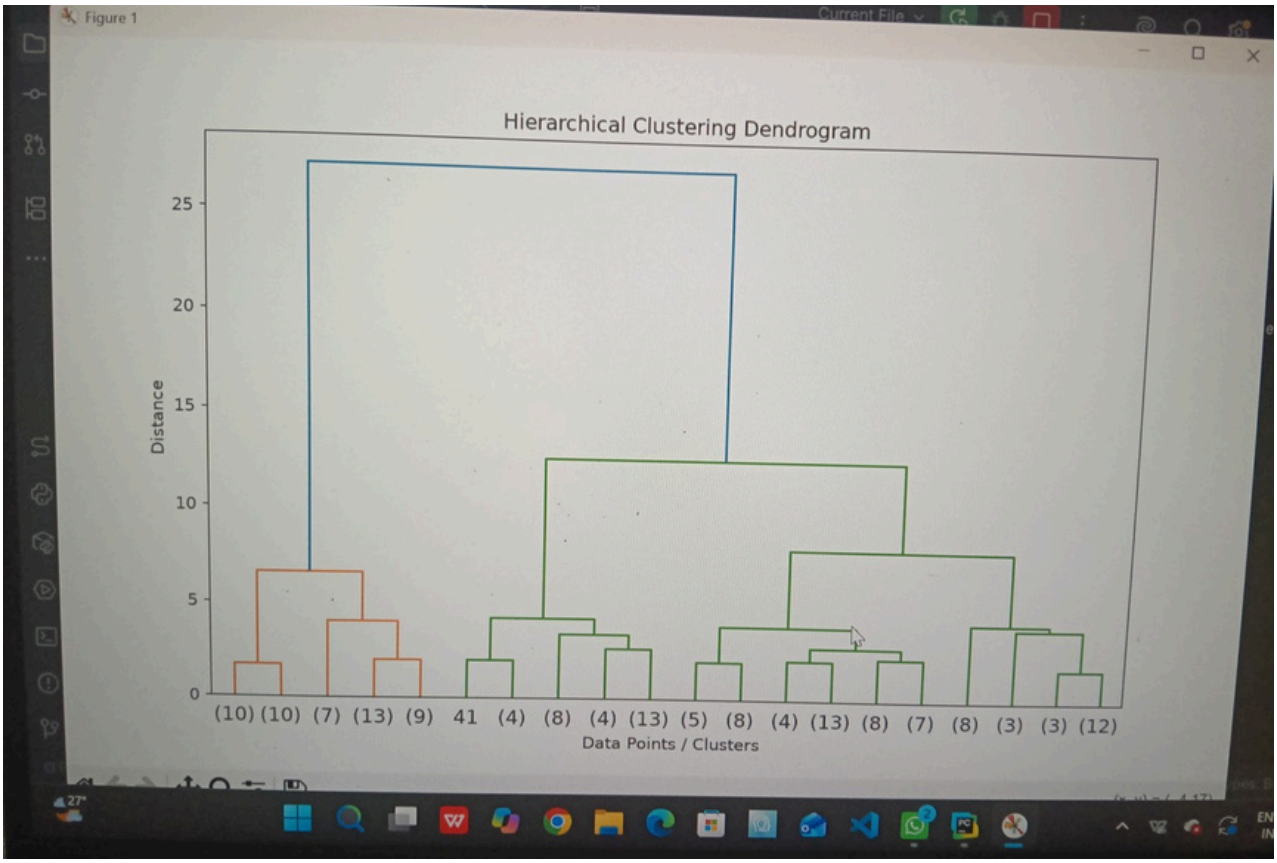
1. Start.
2. Import the required libraries: Pandas, Matplotlib, dendrogram, linkage, and AgglomerativeClustering.
3. Load the dataset from data.csv.
4. Select Hours_Studied and Marks as input features.
5. Create a linkage matrix using the Ward linkage method.
6. Plot the dendrogram to visualize the hierarchical relationships between data points.
7. Create an Agglomerative Clustering model with three clusters.
8. Apply the model to the selected data and predict cluster labels.
9. Add the cluster labels to the dataset as a new column named Cluster.
10. Display the clustered dataset.
11. Plot a scatter plot showing the three clusters using different colors.
12. Display the graph.
13. Stop.






PythonProject14
master
Current File
Unlock Pro
Project Files
C:\Users\tandr\PycharmProjects\Python
.idea
.venv
11259A389(1).py
11259A389(2).py
11259A389(3).py
11259A389(4).py
11259A389(5).py
11259A389.py
datascience.py
HierarchicalClustering.py
ImplementBoosting.py
NeuralNetwork.py
RandomForest.py
ReinforcementLearning.py
SVMClassification.py
Extensions
Scratches
Run
HierarchicalClustering
Dataset Shape: (150, 4)
Number of Samples: 150
Number of Features: 4
Features:
['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width (cm)']
Cluster Labels:
[1 1 1





Result

Thus, the Hierarchical Clustering algorithm was implemented successfully using Ward linkage and Agglomerative Clustering. The hierarchical relationships were visualized using a dendrogram, and the final three clusters were displayed using a scatter plot.

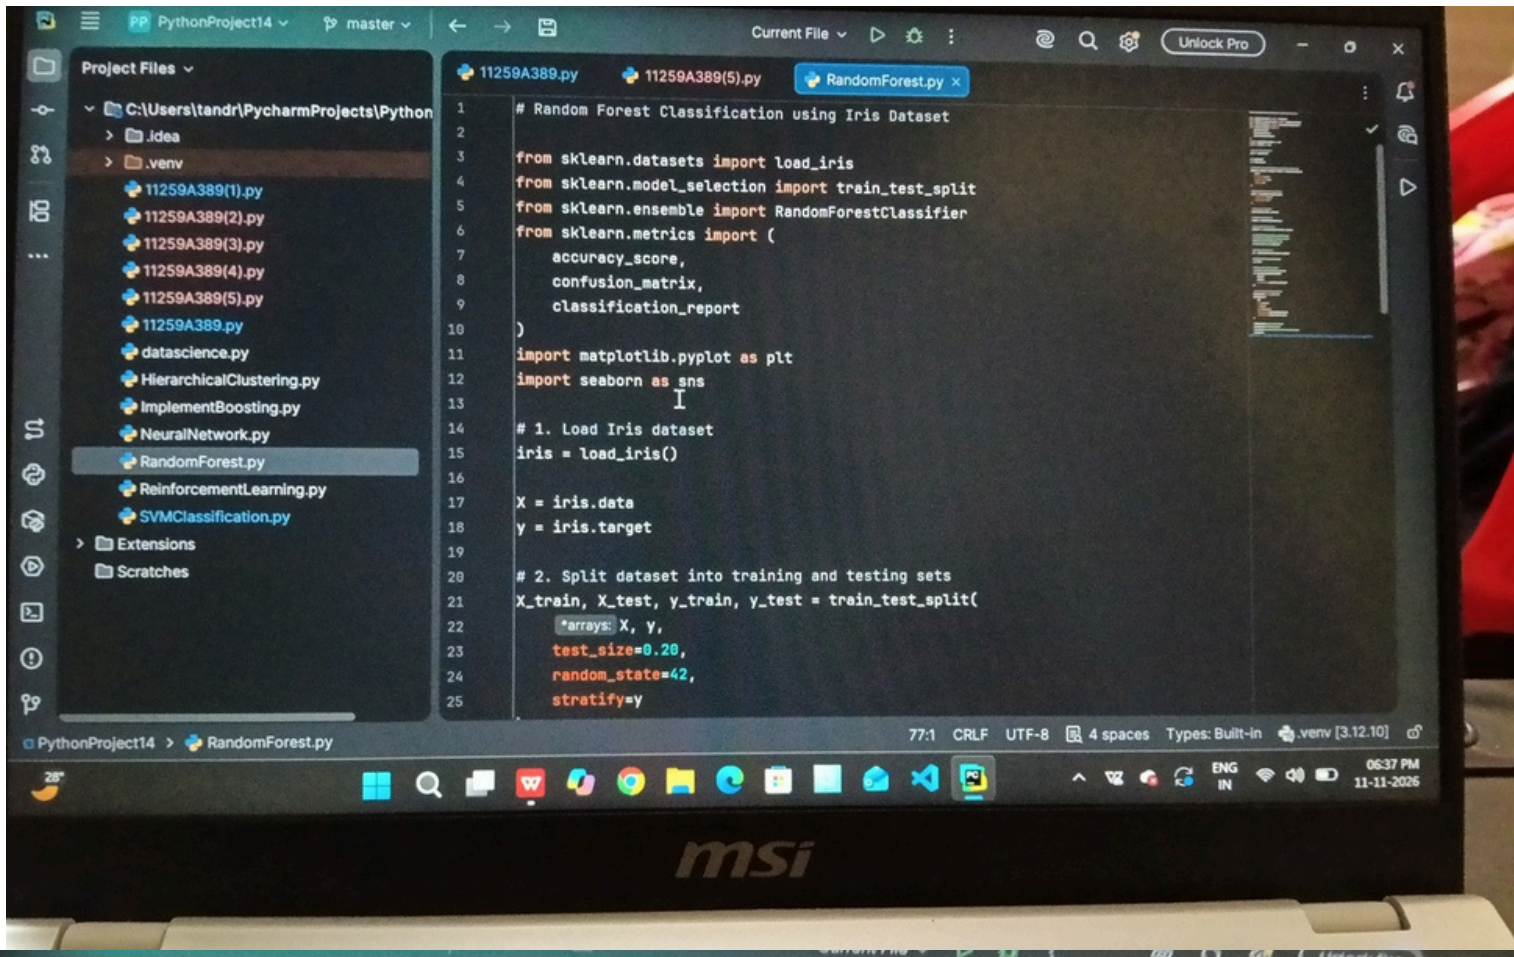
Experiment 7: Implement Random Forest and Analyze Performance

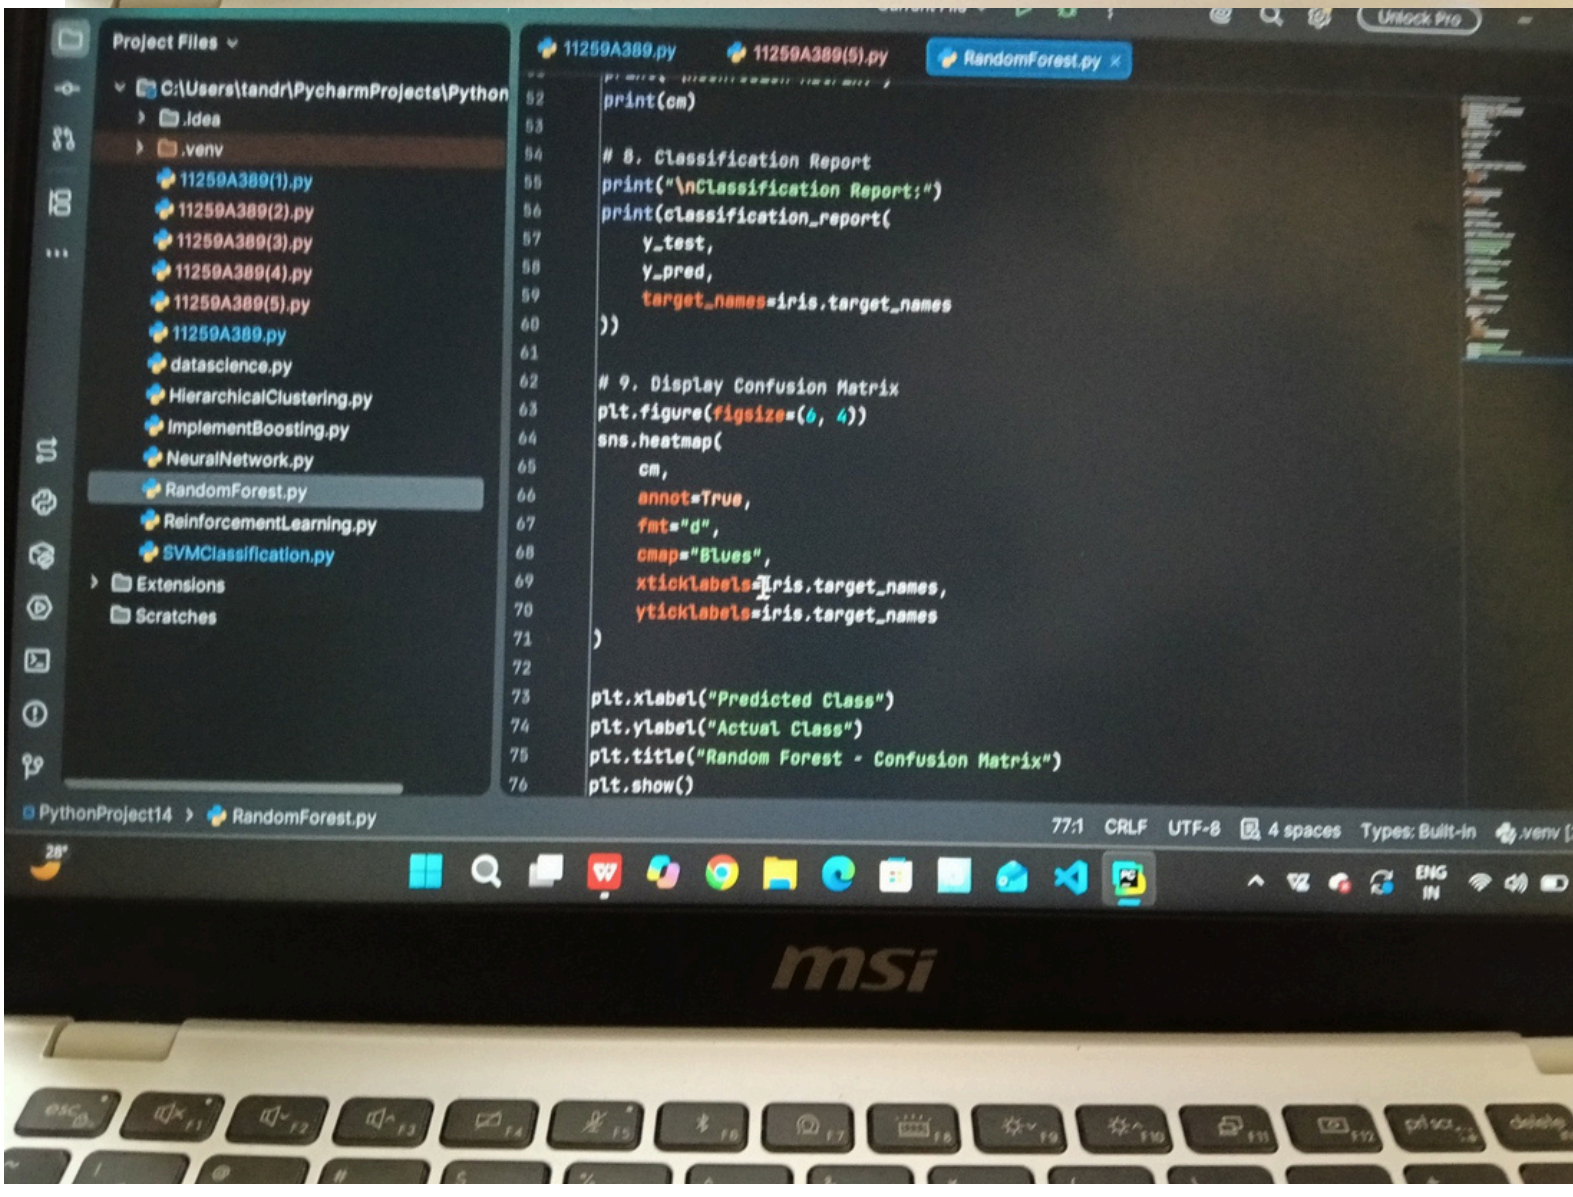
Aim

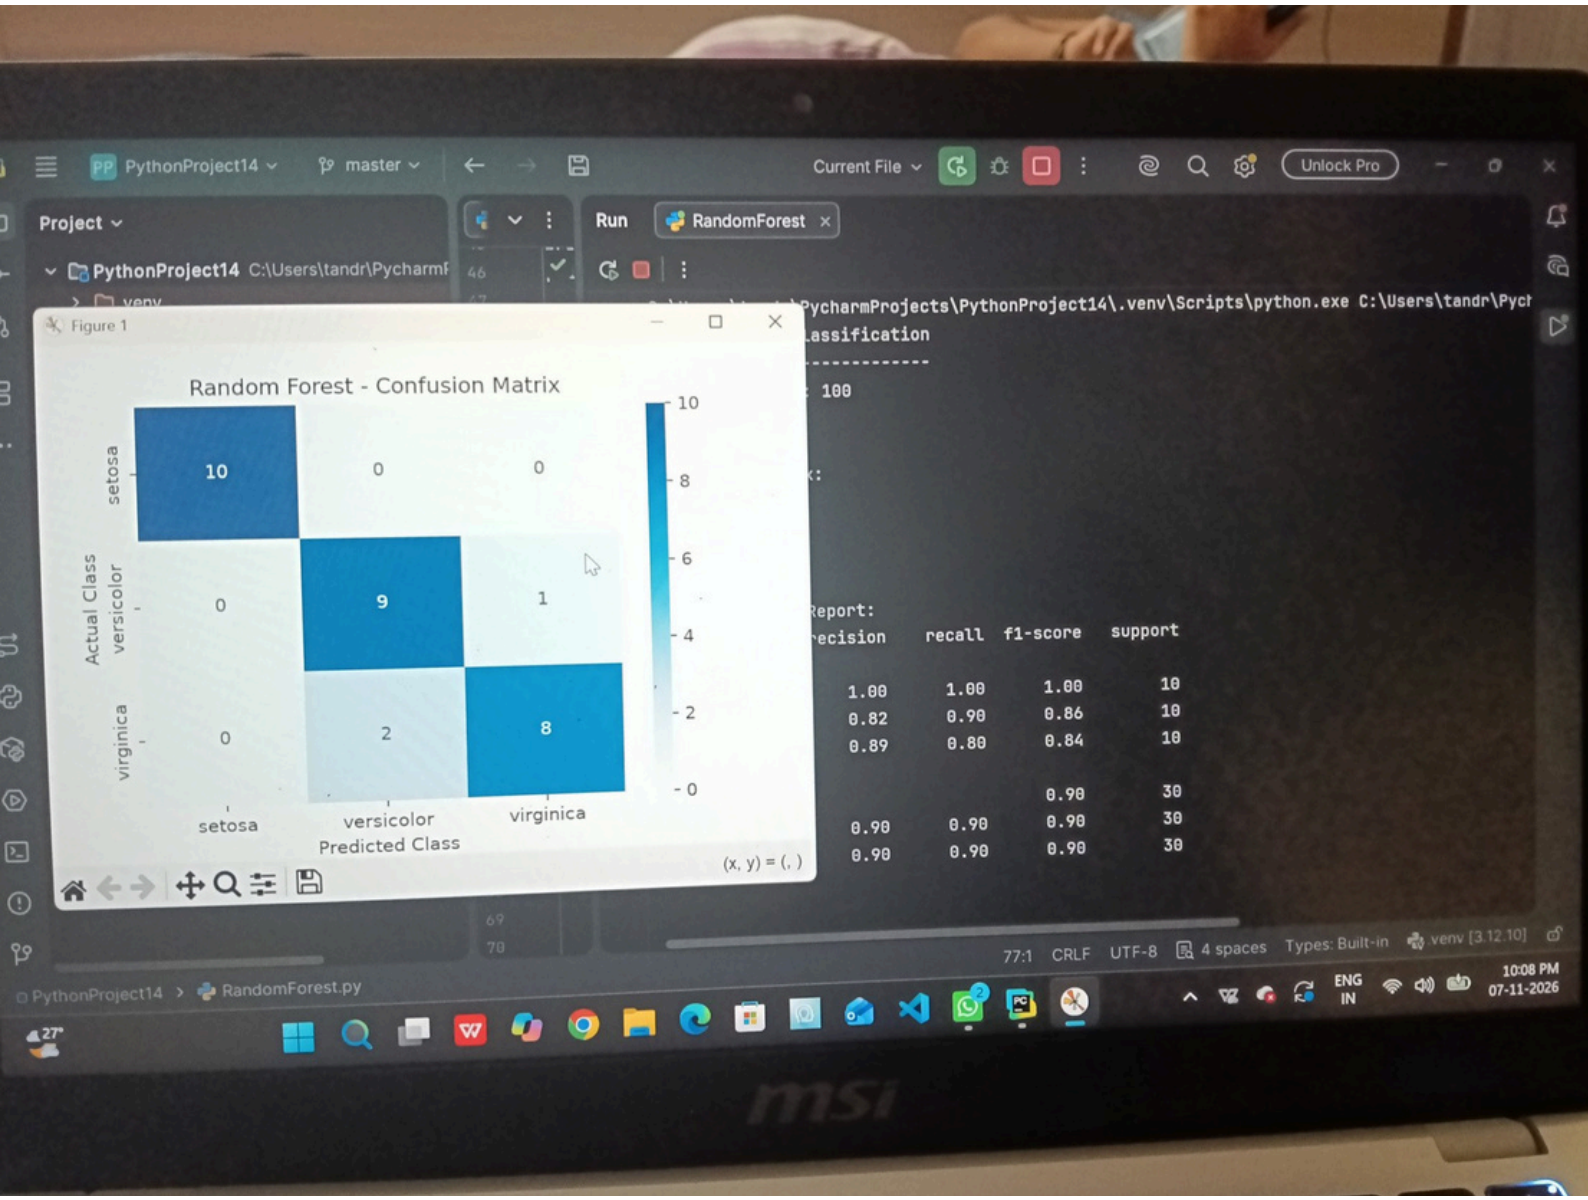
To implement the Random Forest classification algorithm to predict students' results based on hours studied, attendance, and assignments, and analyze its performance using accuracy, confusion matrix, classification report, and feature importance visualization.

Algorithm

1. Start.
2. Import the required libraries: Pandas, train_test_split, RandomForestClassifier, evaluation metrics, and Matplotlib.
3. Load the dataset from data.csv.
4. Select Hours_Studied, Attendance, and Assignments as input features.
5. Select Result as the target variable.
6. Split the dataset into 80% training data and 20% testing data using train_test_split().
7. Create a Random Forest classifier with 100 decision trees and random_state=42.
8. Train the model using the training dataset.
9. Predict the results for the testing dataset.
10. Calculate and display the accuracy score.
11. Display the confusion matrix and classification report.
12. Extract the feature importance values from the trained model.
13. Plot a bar graph showing the importance of each feature.
14. Display the graph.
15. Stop.


PythonProject14
master
Current File
Unlock Pro
Project Files
C:\Users\tandr\PycharmProjects\Python
.idea
.venv
11259A389(1).py
11259A389(2).py
11259A389(3).py
11259A389(4).py
11259A389(5).py
11259A389.py
datascience.py
HierarchicalClustering.py
ImplementBoosting.py
NeuralNetwork.py
RandomForest.py
ReinforcementLearning.py
SVMClassification.py
Extensions
Scratches
11259A389.py
11259A389(5).py
RandomForest.py
Random Forest Classification using Iris Dataset
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import (
accuracy_score,
confusion_matrix,
classification_report
)
import matplotlib.pyplot as plt
import seaborn as sns
1. Load Iris dataset
iris = load_iris()
X = iris.data
y = iris.target
2. Split dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(
*arrays: X, y,
test_size=0.20,
random_state=42,
stratify=y
PythonProject14 > RandomForest.py
77:1 CRLF UTF-8 4 spaces Types: Built-in .venv [3.12.10]
ENG IN
06:37 PM 11-11-2026
msi


Project Files
C:\Users\tandr\PycharmProjects\Python
.idea
.venv
11259A389(1).py
11259A389(2).py
11259A389(3).py
11259A389(4).py
11259A389(5).py
11259A389.py
datascience.py
HierarchicalClustering.py
ImplementBoosting.py
NeuralNetwork.py
RandomForest.py
ReinforcementLearning.py
SVMClassification.py
Extensions
Scratches
11259A389.py
11259A389(5).py
RandomForest.py
print(cm)
8. Classification Report
print("\nClassification Report:")
print(classification_report(
y_test,
y_pred,
target_names=iris.target_names
))
9. Display Confusion Matrix
plt.figure(figsize=(6, 4))
sns.heatmap(
cm,
annot=True,
fmt="d",
cmap="Blues",
xticklabels=iris.target_names,
yticklabels=iris.target_names
)
plt.xlabel("Predicted Class")
plt.ylabel("Actual Class")
plt.title("Random Forest - Confusion Matrix")
plt.show()
PythonProject14 > RandomForest.py
77:1 CRLF UTF-8 4 spaces Types: Built-in
ENG IN
msi



Result

Thus, the Random Forest classification algorithm was implemented successfully to predict students' results. Its performance was analyzed using accuracy, confusion matrix, precision, recall, and F1-score. Feature importance was also visualized using a bar graph.

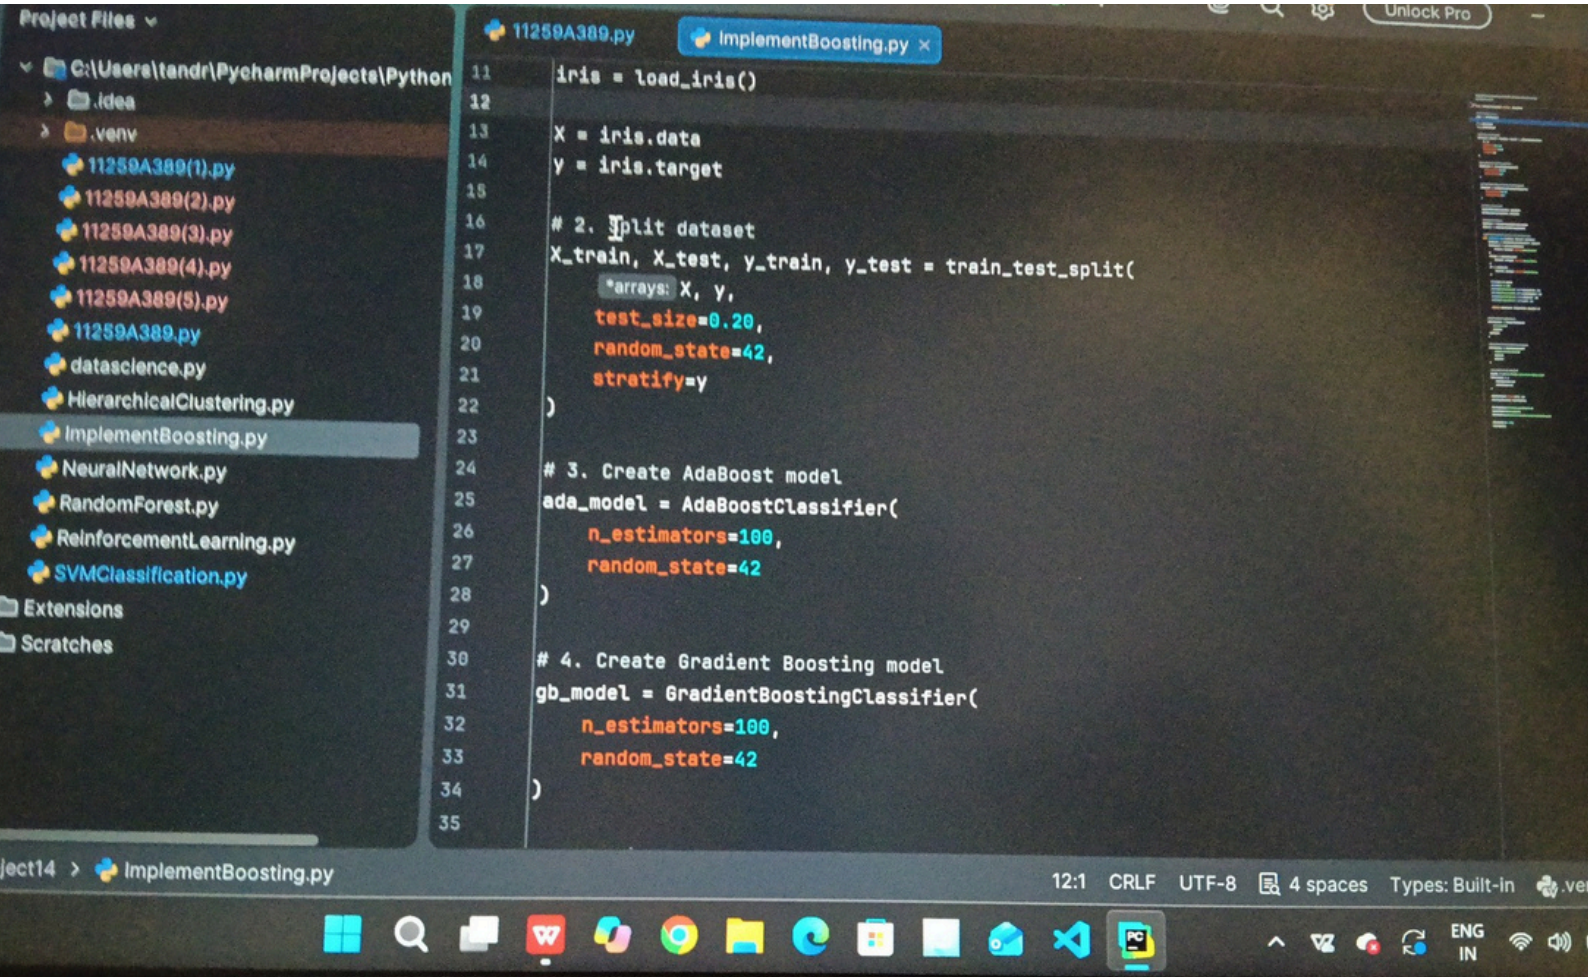
Experiment 8: Implement Boosting Technique (AdaBoost) and Analyze Performance

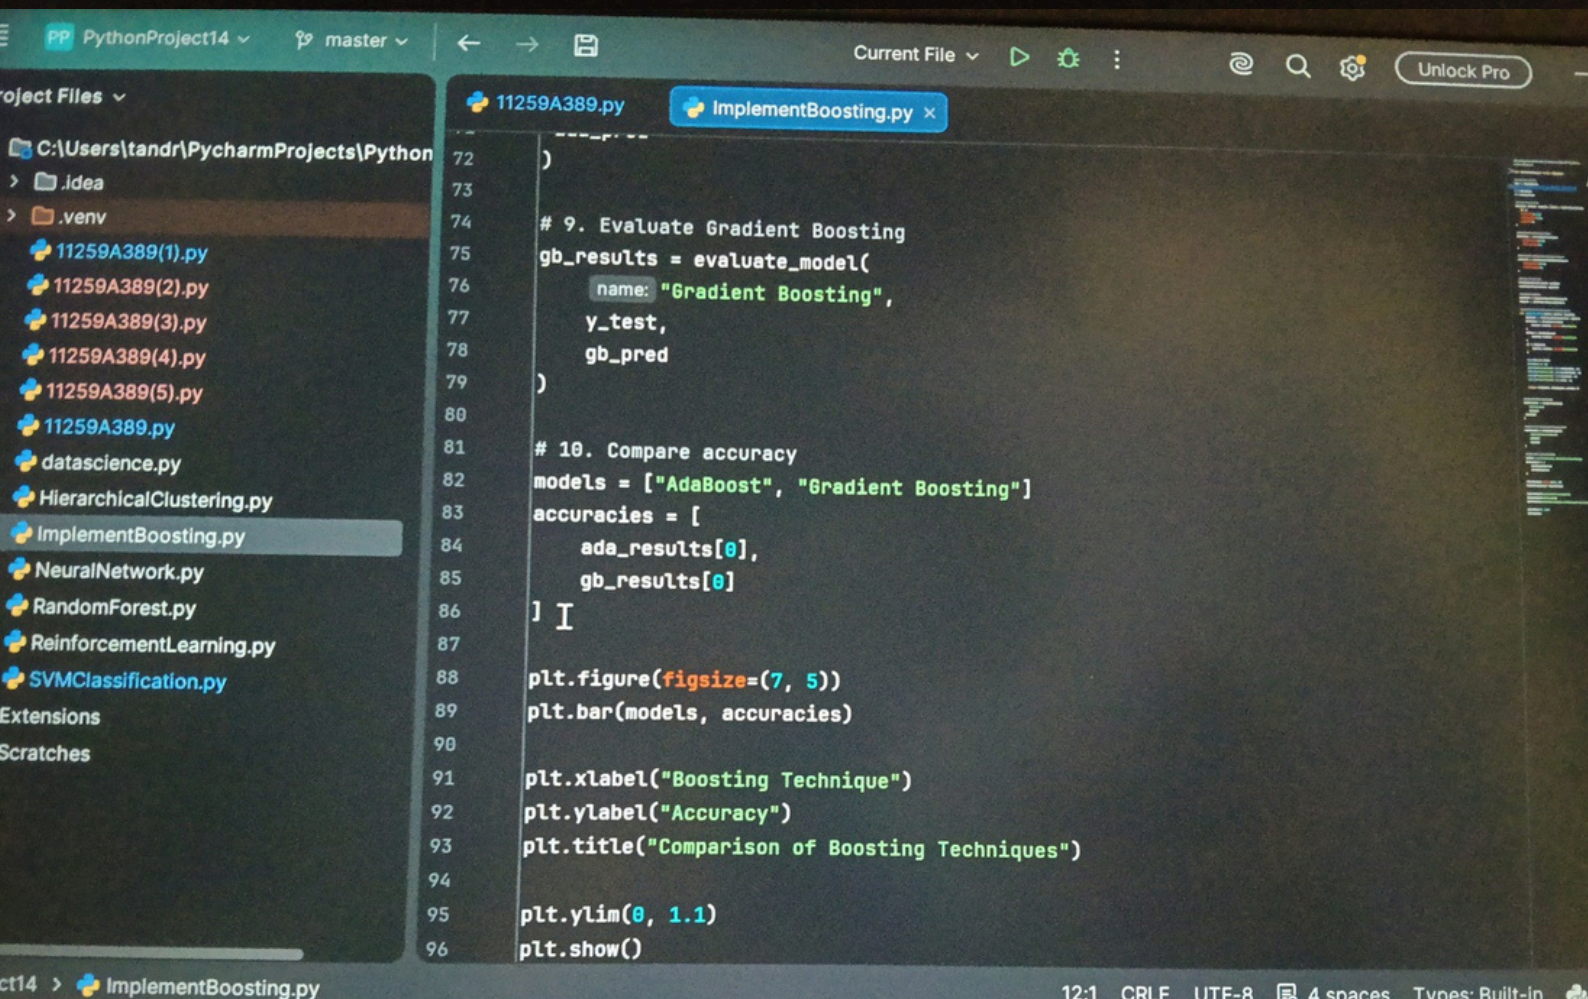
Aim

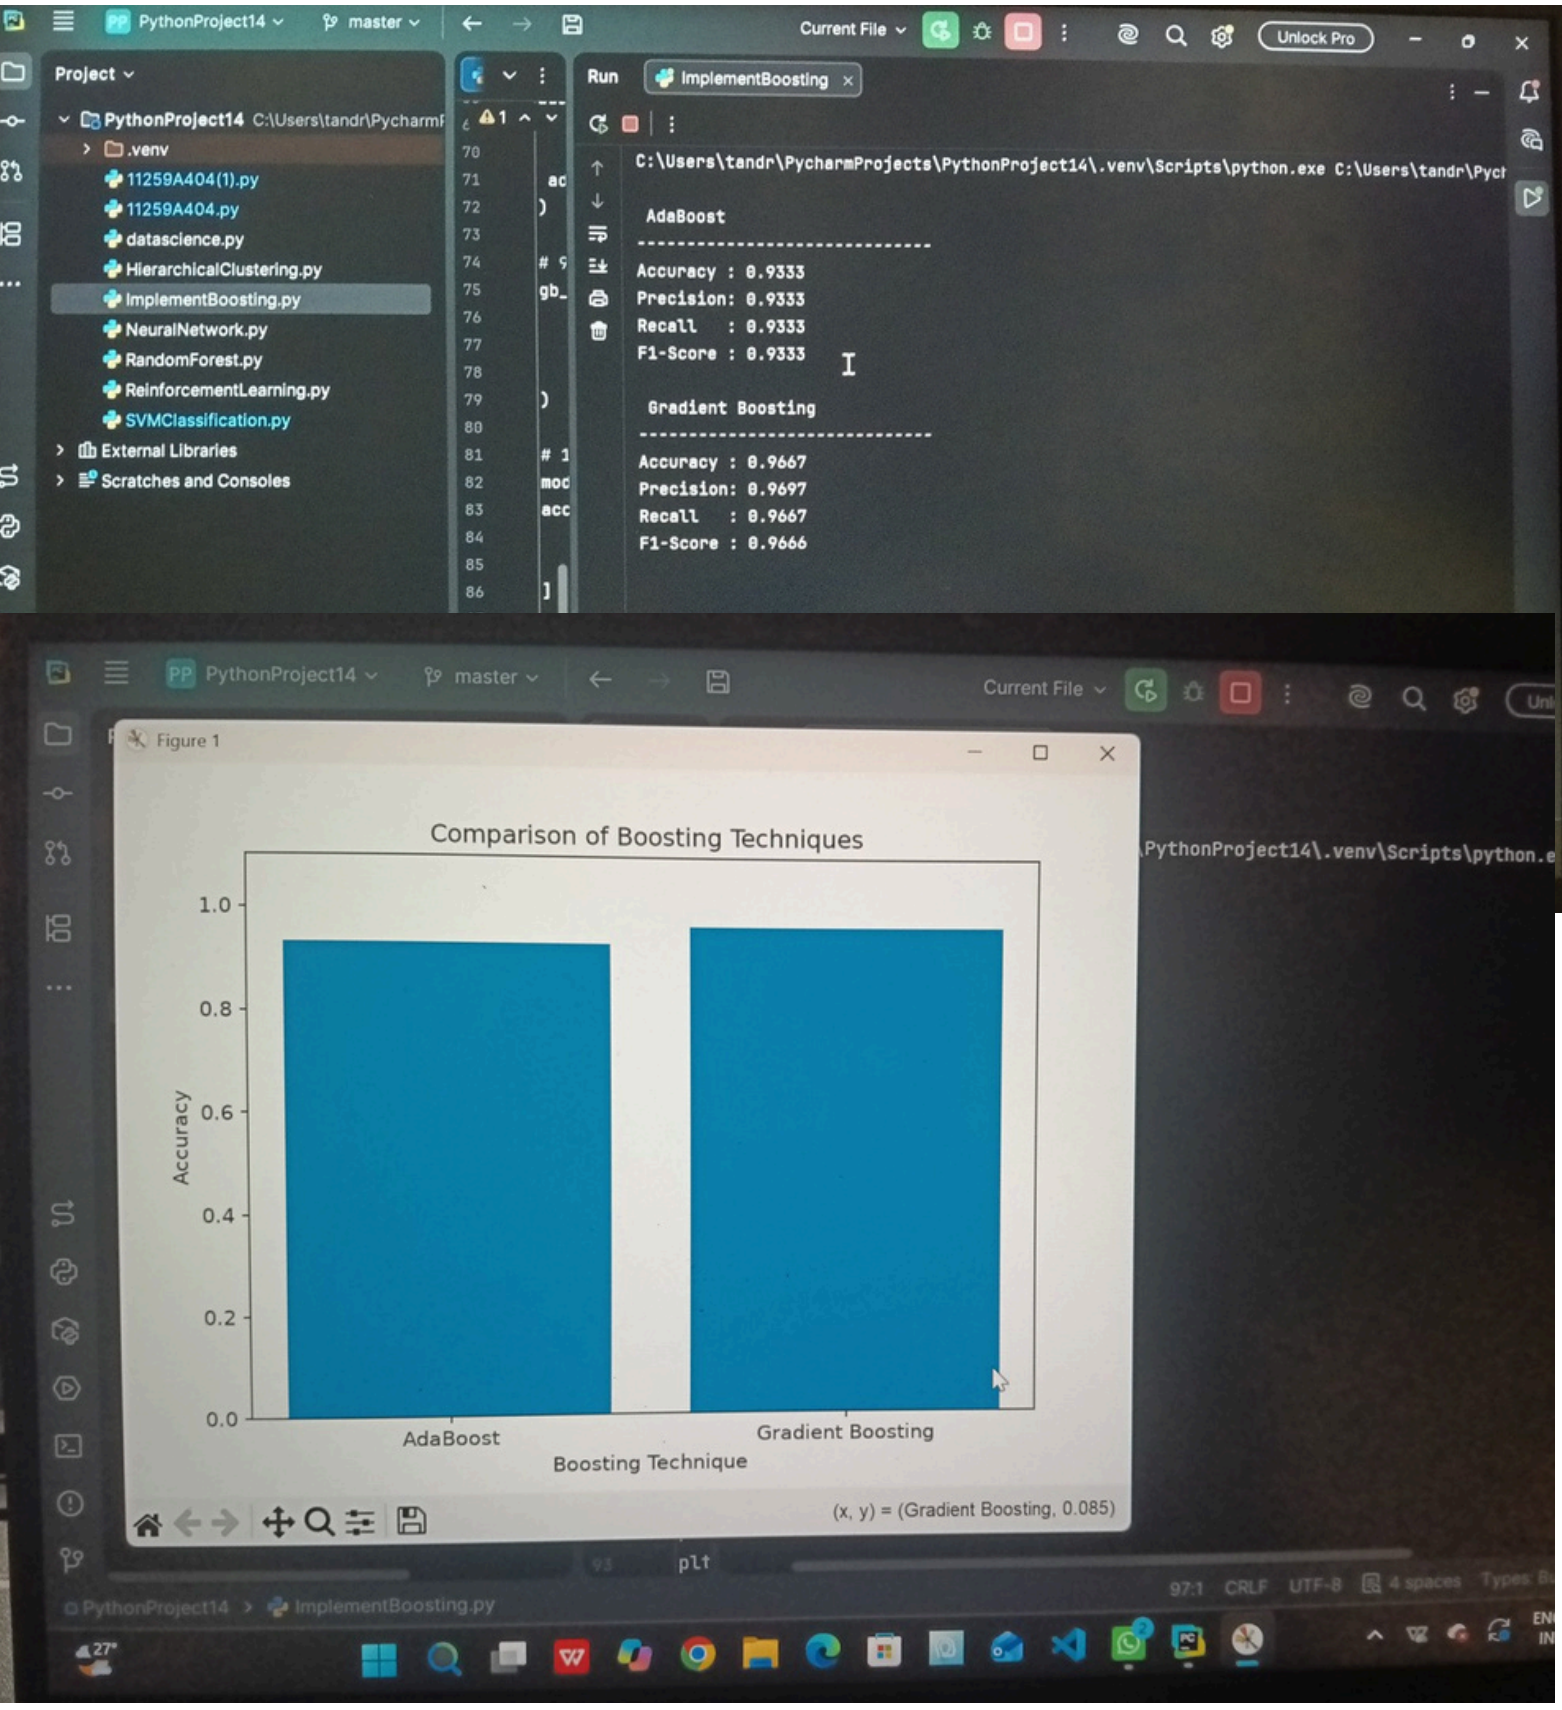
To implement the AdaBoost classification algorithm to predict students' results based on hours studied, attendance, and assignments, and analyze its performance using accuracy, confusion matrix, classification report, and feature importance visualization.

Algorithm

1. Start.
2. Import the required libraries: Pandas, Matplotlib, train_test_split, AdaBoostClassifier, and evaluation metrics.
3. Load the dataset from data.csv.
4. Select Hours_Studied, Attendance, and Assignments as input features.
5. Select Result as the target variable.
6. Split the dataset into 80% training data and 20% testing data.
7. Create an AdaBoost classifier with 50 estimators.
8. Train the model using the training dataset.
9. Predict the results for the testing dataset.
10. Calculate and display the accuracy score.
11. Display the confusion matrix and classification report.
12. Extract the feature importance values.
13. Plot a bar graph showing the importance of each feature.
14. Display the graph.
15. Stop


Project Files
C:\Users\tandr\PycharmProjects\Python
.idea
.venv
11259A389(1).py
11259A389(2).py
11259A389(3).py
11259A389(4).py
11259A389(5).py
11259A389.py
datascience.py
HierarchicalClustering.py
ImplementBoosting.py
NeuralNetwork.py
RandomForest.py
ReinforcementLearning.py
SVMClassification.py
Extensions
Scratches
11259A389.py
ImplementBoosting.py
Unlock Pro
iris = load_iris()
X = iris.data
y = iris.target
2. Split dataset
X_train, X_test, y_train, y_test = train_test_split(
*arrays: X, y,
test_size=0.20,
random_state=42,
stratify=y
)
3. Create AdaBoost model
ada_model = AdaBoostClassifier(
n_estimators=100,
random_state=42
)
4. Create Gradient Boosting model
gb_model = GradientBoostingClassifier(
n_estimators=100,
random_state=42
)
ject14 > ImplementBoosting.py
12:1 CRLF UTF-8 4 spaces Types: Built-in
ENG
IN


PythonProject14
master
Current File
Unlock Pro
oject Files
C:\Users\tandr\PycharmProjects\Python
.idea
.venv
11259A389(1).py
11259A389(2).py
11259A389(3).py
11259A389(4).py
11259A389(5).py
11259A389.py
datascience.py
HierarchicalClustering.py
ImplementBoosting.py
NeuralNetwork.py
RandomForest.py
ReinforcementLearning.py
SVMClassification.py
Extensions
Scratches
11259A389.py
ImplementBoosting.py
)
9. Evaluate Gradient Boosting
gb_results = evaluate_model(
name: "Gradient Boosting",
y_test,
gb_pred
)
10. Compare accuracy
models = ["AdaBoost", "Gradient Boosting"]
accuracies = [
ada_results[0],
gb_results[0]
]
plt.figure(figsize=(7, 5))
plt.bar(models, accuracies)
plt.xlabel("Boosting Technique")
plt.ylabel("Accuracy")
plt.title("Comparison of Boosting Techniques")
plt.ylim(0, 1.1)
plt.show()
ct14 > ImplementBoosting.py
12:1 CRLF UTF-8 4 spaces Types: Built-in
ENG
IN



Result

Thus, the AdaBoost classification algorithm was implemented successfully. Its performance was analyzed using accuracy, confusion matrix, precision, recall, and F1-score, and feature importance was visualized using a bar graph.

Experiment 9: Implement a Basic Reinforcement Learning Model (Q-Learning)

Aim

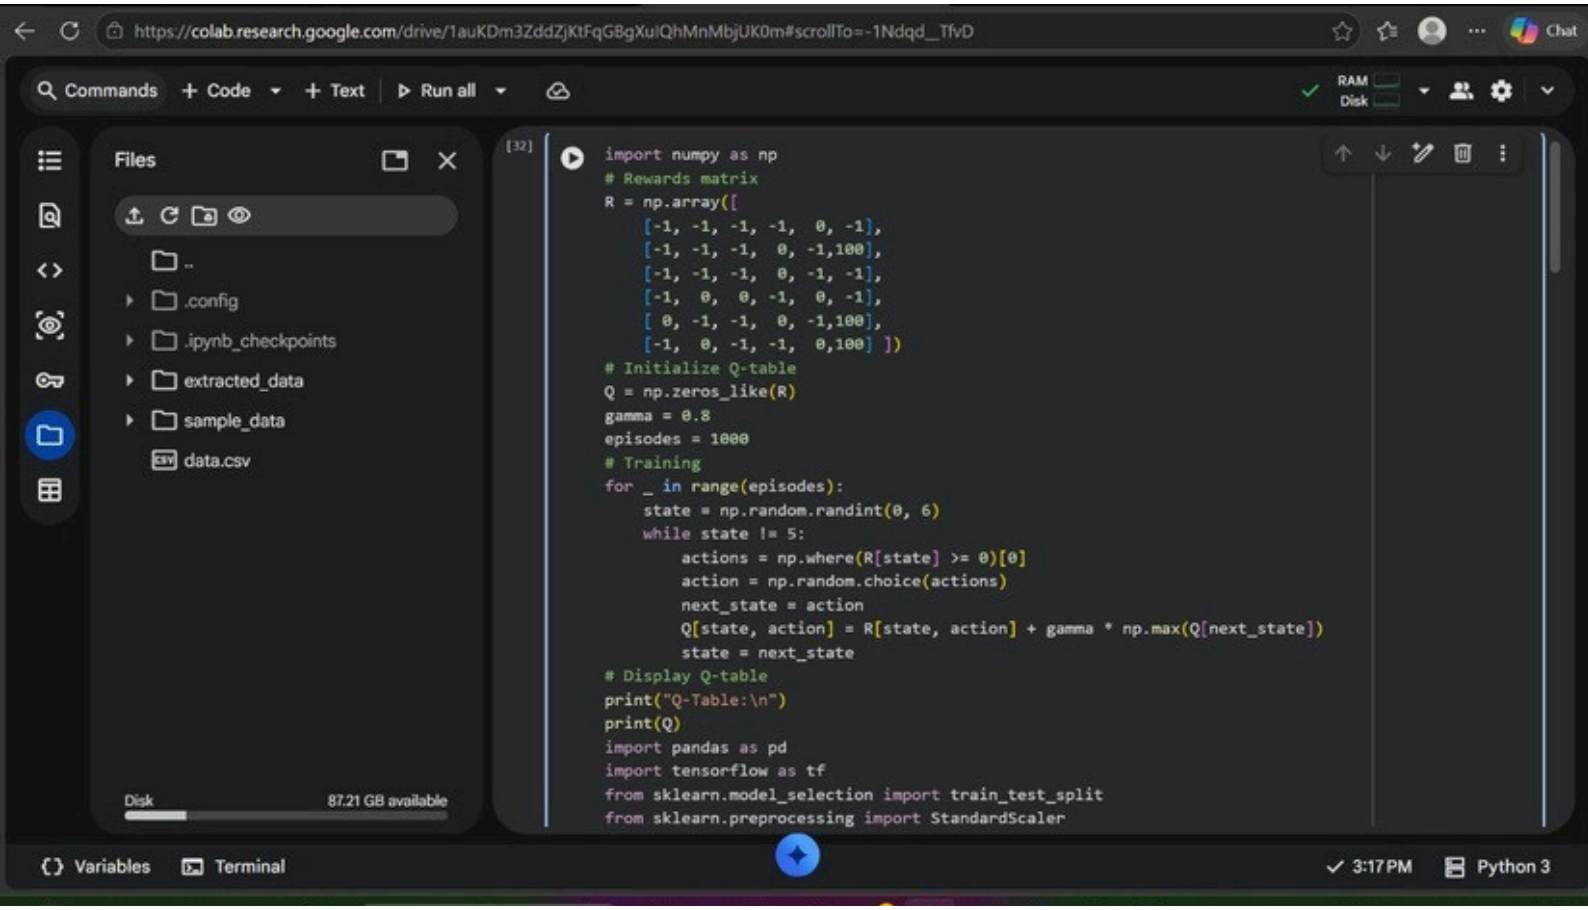
To implement a basic reinforcement learning model using the Q-Learning algorithm to learn the optimal actions for reaching a target state using a reward matrix and a Q-table.

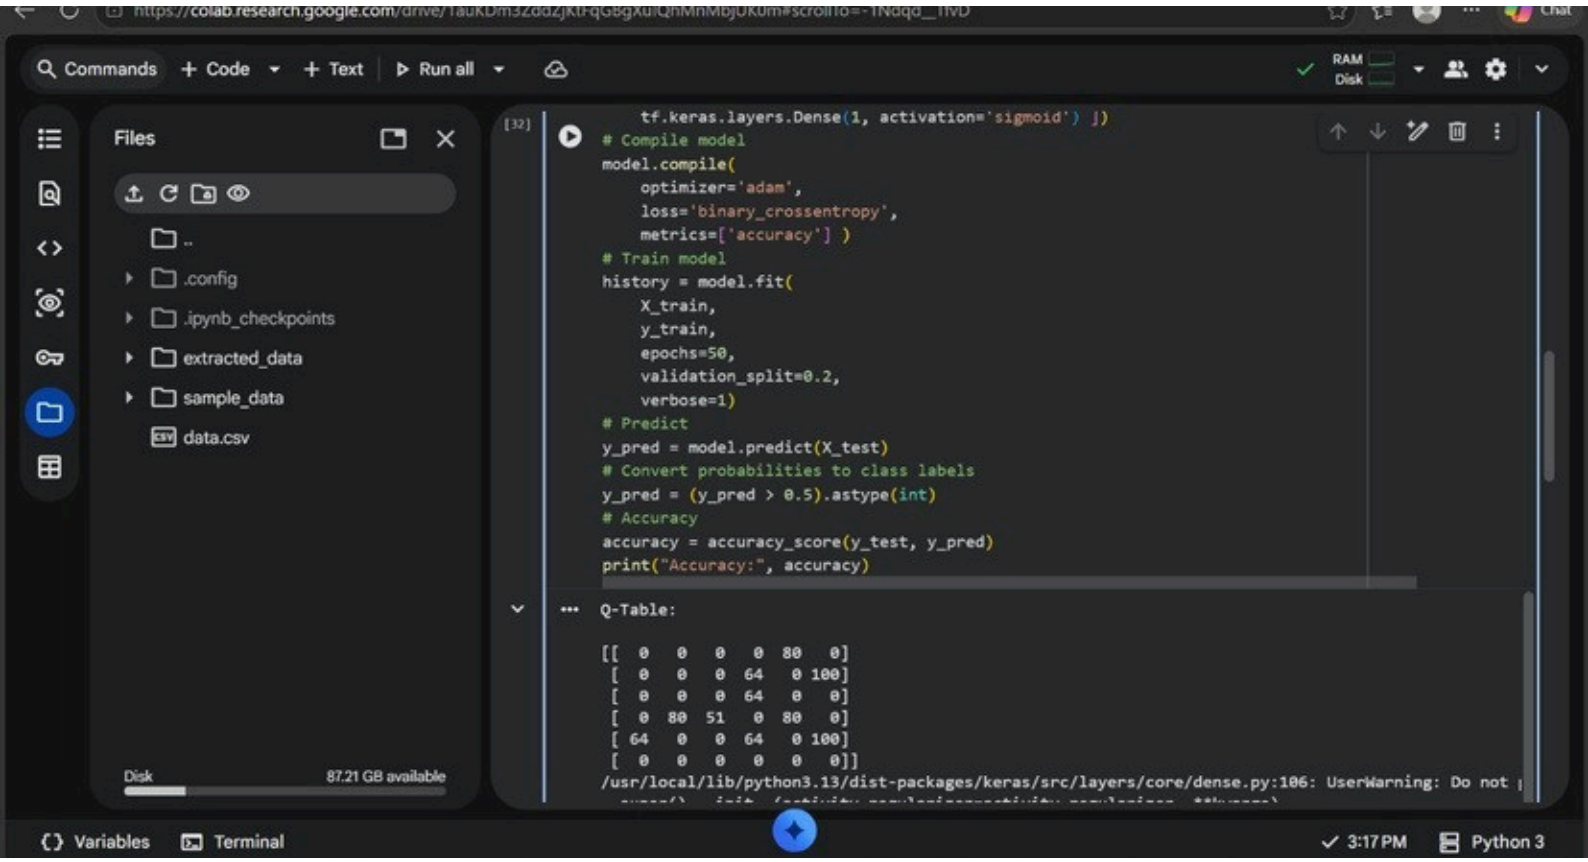
Alg1**o**. **r**S**i**ta**h**r**m**t.

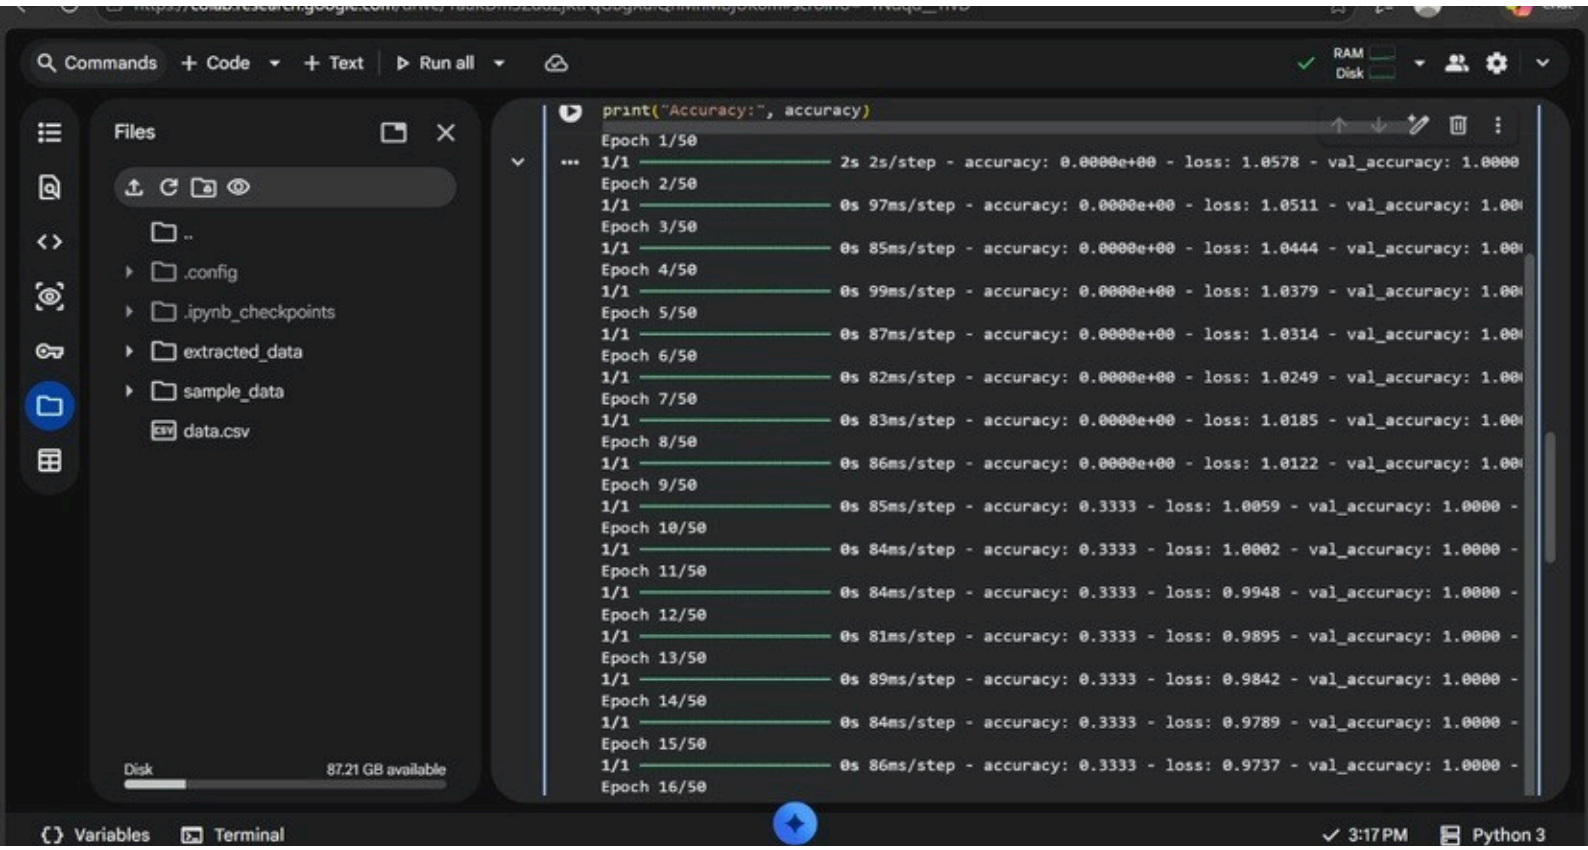
2. Import the NumPy library.
3. Define the reward matrix representing the rewards for moving between states.
4. Initialize the Q-table with zeros.
5. Set the discount factor \(\gamma = 0.8\) and the number of training episodes to 1000.
6. Randomly select an initial state.
7. Identify the possible actions from the current state using the reward matrix.
8. Randomly select one of the available actions.
9. Determine the next state based on the selected action.
10. Update the Q-table using the Q-Learning update equation:

\[Q(s,a)=R(s,a)+\gamma\max_{a'}Q(s',a') \]

11. Continue until the target state (state 5) is reached.
12. Repeat the training process for 1000 episodes.
13. Display the final Q-table.
14. Stop


https://colab.research.google.com/drive/1auKDm3ZddZjKtFqGBgXulQhMnMbjUK0m#scrollTo=-1Ndqd_TfvD
Commands
+ Code
+ Text
Run all
RAM
Disk
Files
..
.config
.ipynb_checkpoints
extracted_data
sample_data
data.csv
Disk
87.21 GB available
[32]
import numpy as np
Rewards matrix
R = np.array([
[-1, -1, -1, -1, 0, -1],
[-1, -1, -1, 0, -1,100],
[-1, -1, -1, 0, -1, -1],
[-1, 0, 0, -1, 0, -1],
[0, -1, -1, 0, -1,100],
[-1, 0, -1, -1, 0,100]])
Initialize Q-table
Q = np.zeros_like(R)
gamma = 0.8
episodes = 1000
Training
for _ in range(episodes):
state = np.random.randint(0, 6)
while state != 5:
actions = np.where(R[state] >= 0)[0]
action = np.random.choice(actions)
next_state = action
Q[state, action] = R[state, action] + gamma * np.max(Q[next_state])
state = next_state
Display Q-table
print("Q-Table:\n")
print(Q)
import pandas as pd
import tensorflow as tf
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
Variables
Terminal
3:17 PM
Python 3


Commands
+ Code
+ Text
Run all
RAM
Disk
Files
..
.config
.ipynb_checkpoints
extracted_data
sample_data
data.csv
Disk
87.21 GB available
[32]
tf.keras.layers.Dense(1, activation='sigmoid')])
Compile model
model.compile(
optimizer='adam',
loss='binary_crossentropy',
metrics=['accuracy'])
Train model
history = model.fit(
X_train,
y_train,
epochs=50,
validation_split=0.2,
verbose=1)
Predict
y_pred = model.predict(X_test)
Convert probabilities to class labels
y_pred = (y_pred > 0.5).astype(int)
Accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
Q-Table:
[[0 0 0 0 80 0]
[0 0 0 64 0 100]
[0 0 0 64 0 0]
[0 80 51 0 80 0]
[64 0 0 64 0 100]
[0 0 0 0 0 0]]
/usr/local/lib/python3.13/dist-packages/keras/src/layers/core/dense.py:106: UserWarning: Do not
Variables
Terminal
3:17 PM
Python 3



Result

Thus, the Q-Learning reinforcement learning model was implemented successfully. The agent learned state-action values through repeated interaction with the reward matrix, and the learned values were displayed in the Q-table.